

Unethically scraping r/transtimelines posts to inaccurately predict gender transition outcomes

Anonymous

April 1, 2025

Abstract

This report details a system for the pressing problem of predicting how a person, who is currently a guy, would look like as a girl, realistically. The system does not actually really work and is just a prompt for Gemini along with a dubiously ethical evaluation dataset. However, we got too far to just abandon this project, so let this ship sink here.

1 Prior art

Yudkowsky [2009] has considered the problem of gender transition and predicting how it will go in the rationalist philosophical framework and concluded that the idea is silly and infeasible for the same reasons AI value alignment is. However, it primarily focuses on changing the *brain* of a person to be female, not the body and face, and the reason for imagining an idealized magical transition is typically trying to escape from the fear of not looking good after transition. Real-life transition is possible, and there is a distribution on outcomes conditional on pre-existing conditions.

FaceApp [Wikipedia, 2016] is an early “AI”-based mobile app for editing photos of faces. It has a gender filter feature that can turn a selected face to like like that of the opposite sex. According [Wikipedia, 2016] to Wikipedia [Wikipedia, 2001], the “...gender change transformations of FaceApp have attracted particular interest from the LGBT and transgender communities, due to their ability to realistically simulate the appearance of a person as the opposite gender.” According to one of Wikipedia’s primary sources [Blank, 2021],¹ the app is inadequate not because it is inaccurate,² but because it erases the nonbinariness and fluidity of gender identity and expression or someting. While this issue is not relevant to the topic of this paper, we have previously touched upon it and are happy to self-cite [Anonymous, 2024].

2 Data

The r/transtimelines subreddit [Reddit Community, n.d.] is an uplifting subreddit to hopescroll through. Many transgenders share their stories visually and inspire others to take the plunge. And that is good. But sharing photos on the internet is dangerous.⁴⁵ In this paper, we will treat these brave individuals as unconsenting data donors for our devious little experiments, the results of which we will only share as summary statistics.

We found all image posts on the subreddit as of Fall 2024 through unknown means.⁶ We then filtered for imgur, reddit and tumblr direct image links.⁸ We downloaded them with 4 processes running a **certain program** from behind a VPN.⁹ We filtered again with a multi-step pipeline:

¹Which itself cites some Twitter polls.

²Accuracy is supposedly one of FaceApp’s selling points according to the CEO [Goncharov, 2016].³

³No, that reference isn’t clickable ☹

⁴Opsec, citation retracted.

⁵Especially if you are transgender. Because whatever you share online could be used to deanonymize and harass you by thousands of enthusiasts. So, by scraping this data we are likening ourselves to these abusers. Disgusting.

⁶The authors tried to reconstruct the method used to obtain `posts.json` file with the URLs of the posts. The most likely explanation is that we simply downloaded it from [REDACTED].⁷

⁷Those who know, know. Others shouldn’t try to replicate a paper the authors explicitly express regret writing.

⁸This filtering step actually failed and we downloaded *all* images posted to this site. There was a `redgifs` video. I was in public, and I mindlessly clicked on the downloads directory and opened a file. It was about what you would expect.

⁹Or maybe not. Maybe I’m on some watchlist as a weirdo downloading images of trans people.

1. Look for really prominent vertical seams in a Canny edge-detected version of the image. In the images we’re looking for, there’s always exactly one vertical seam with a man on the left and a woman on the right.¹⁰ We used some threshold to discard images without a prominent seam.¹¹
2. Detect bounding boxes of people in the image with MobileNetSSD. Check that there’s two and that they roughly overlap with the two halves of the image.
3. Trim out rows and columns of pure white space on the borders of the image.
4. Send the image to a Gemini 1.5 Flash instance with the output format of¹³

```
{"isNotConfusing": bool, "isCollage": bool, "twoPicturesSideBySide": bool, "hasExactlyTwoPeople": bool}
```

This filters out FtM samples.¹⁴¹⁵¹⁷

5. If the image passes all of the vibe checks, cut it in half and save it.

We make the raw and processed data available at [from sigbovik, 2025].* We have about 60k raw images and 1.7k processed image pairs.¹⁸

3 Method

We finetuned a Flux.1 dev lora on side-by-side before-and-after pictures with a rank of 128 and a very high alpha. We used it for inpainting the right part of the picture from the left conditionally. It worked, but was bad.¹⁹

With Gemini Flash 2.0’s native image generation [Kampf, 2025], it would be possible to generate many awesome memes for this paper.²⁰²¹ Oh, and to accomplish the goal of this paper without any training. This paper was originally going to be about training models like Flux et al, but calling APIs seemed more fun than waiting for runs to finish.²²

We initially faced some challenges. The model would refuse to generate transition photos if we said the picture was found online, or presented as belonging to another person in any way. We needed to do a bit of social engineering and pretend the photo belonged to ourselves.²³

¹⁰Philip [2021], but the other way around.

¹¹If you are replicating this, we want to make you suffer and spend... 3 minutes tuning this threshold. Ok, this is pointless.¹²

¹²We looked up the threshold. It’s pretty contingent on our exact setup, but we still won’t reveal it. For decision theory reasons [Ziz, 2017].

¹³The other two conditions are our secret sauce. We know you won’t look at the PDF source code because this is a stupid parody paper and you don’t have any reason to.

¹⁴Pretend you didn’t see that.

¹⁵Secret sauce aside, is this transmasculine erasure? Yes, kind of. Except... This is a personal project. I have no use for these photos, they are from a completely different distribution than the one I need for my “usecase”. I don’t want to forward these images to Google’s servers more times than necessary because it would be completely pointless.¹⁶

¹⁷But isn’t this entire project pointless? No, it’s the MVP I need to get out to satisfy the obsession with submitting a paper this exact time of the year as a form of displacement.

*It is gated by the author’s approval because we don’t want random internet people to be able to download 50 thousand people’s deeply personal timeline photos easily.**

We probably won’t let anyone actually access it if we don’t know you personally because we feel bad about scraping the subreddit.*

*** If we scraped your data, I’m sorry and I promise to delete it if Huggingface ever finds a way to do something with all of the data people upload for once instead of just deleting everything.****

**** Is Huggingface better than tech companies at this? At filtering and making use of personal data users give out? That’s a good SIGBOVIK paper topic. Better than this one. This paper is an academic self-harm method.

¹⁸Bottlenecked on API rate limits.

¹⁹It did learn the room layout of a typical redditor! And made the faces look less perfect. When used outside of inpainting mode, it generally just made both faces round and feminine. Which is kind of affirming.

²⁰See footnote 10 for an example.

²¹Although OpenAI’s as of now recent release overshadowed that. And we don’t want to use any images in the main body for Reasons.

²²And we got hooked on generating fake timelines of our own face. It’s kind of unsettling, the image looks like it belongs to a sibling or doppelgänger, and yet there’s a lot of the variation in the images and features because of the flexibility of interpretation of the original image. That is, when it isn’t completely off the mark.

²³This seems harmless, but it felt terrible. I stole someone’s face and was pretending to be them for some useless paper. I decided to use my own photos to reduce guilt. See the footnote above.




Figure 1: Main figure. Only figure. Heroic figure. Human evaluation scores for our Gemini-based pipeline, based on painstakingly collected data.

Once we got past that problem, we needed to emphasize that the transition should be MtF, that the person had to retain their identity and that the transition should be “realistic but optimistic”. We share the full prompt in Appendix A.

We tried to add few-shot examples, but found that they only worsened the realism and identity cohesion. Thus, our method is fully data-free,²⁴

We need to evaluate if our method:

1. Accurately captures the original face.
2. Is more or less conventionally attractive than the post-transition face, and to what degree.

Why these two? The first is a good metric, but the model could simply always leave the face unchanged and claim it achieved perfectly faithful transition. We need to match the distribution to real-life outcomes as an importance sampling scheme that shows outcomes that both fall short and exceed what they will actually see in real life to maximally motivate them. We rank the first on a scale of 1-10,²⁵ and the second also on a scale of 1-10, with 5 representing similar attractiveness to the ground truth picture.

Figure 1 illustrates our results. We generated 307 images²⁶ based on random samples from our dataset [from sigbovik, 2025] and ran them through the human evaluation pipeline.²⁷²⁸ These results show a few interesting patterns:

²⁴And we did all that work collecting the data for nothing.

²⁵Effectively 5-10 because the faces are superficially similar in race and lighting. A score of 1 would represent the opposite of a given face.

²⁶We tried to generate 10 candidates for each of 100 base images, but Gemini has a high failure rate on these prompts and has strict rate limits, so we only got 93 done at a $\approx 30\%$ yield.

²⁷It’s 5AM. My eyes hurt. I can’t tell if I should assign this picture of a bald guy a score of 0 or -1.

²⁸You may notice that the number of images on each tile only sum up to 88. This is because I got tired and decided this was good enough.

- There are a lot of faces with a passability ratio of 0 which are just the original photo or men with short hair.
- A passability ratio of 5 is the sweet spot, and it seems to be easily achieved. The method is roughly calibrated, but sometimes makes the images slightly less pretty.
- Evaluating closeness/realism is hard. It would be easier if I could see the skulls of both of the generations. See Section 4.

4 Conclusion

This is a tech report. It doesn't need a conclusion.²⁹ Future work could include using better models like the new GPT 4o (make up the citation yourself), or comparing to Old World transition predictors such as FaceApp. We could also extend the metrics to use more specialized models. For example, we could add face embedding models to compare pre- and post- transition faces. Or landmark models. Like in the crypto chad app. Incel tech [Hamilton, 2024]. We could use them combined with traditional machine learning to learn the one true metric for passability. We could incorporate the user feedback from the evaluation pipeline.³⁰

To evaluate generative methods for this task, we need only a measure of closeness to the original face and not a passability metric. As long as the closeness is maximized, we can sample along the manifold of conventional and unconventional attractiveness and receive faithful possible snapshots. And if the results are not what you expected? Well, they won't be in real life. If they will be.

5 Acknowledgements

Contrary to what our use of first person plural pronouns would lead you to believe, we are actually one person. We just like to use "we" because it sounds cool and makes us feel less lonely.

The computations and scraping ran mostly on our computer, and partially in Google Colaboratory. We used the Gemini API's free daily credits for this. This is very valuable novel research Google must be happy to see it enter the world.

References

- Anonymous. Gender is complex: pulling the laplacian eigengender from relationship graphs. In *Proceedings of SIGBOVIK 2024: The Eighteenth Annual Intercalary Robot Dance Party in Celebration of Workshop on Symposium About 26th Birthdays; In Particular, That of Harry Q. Bovik*, April 2024.
- A. Blank. Faceapp's transgender filter sparks controversy. *Daily Dot*, 2021. URL <https://www.dailydot.com/irl/faceapp-transgender/>.
- A. from sigbovik. Reddit transtimelines dataset, 2025. URL <https://huggingface.co/datasets/sigbovik/reddit-transtimelines>.
- P. Hamilton. Incel.tech, 2024. URL <https://knowyourmeme.com/memes/sites/incel-tech>.
- K. Kampf. Experiment with gemini 2.0 flash native image generation, March 12 2025. URL <https://developers.googleblog.com/en/experiment-with-gemini-20-flash-native-image-generation/>.
- Philip. Dramatic crossroads, 2021. URL <https://knowyourmeme.com/photos/2046343-dramatic-crossroads>.
- Reddit Community. r/transtimelines, n.d. URL <https://www.reddit.com/r/transtimelines/>.
- Wikipedia. Wikipedia — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=Wikipedia&oldid=1282780745>, 2001.

²⁹I swear I'm an academic.

³⁰We could do a lot of things, but we didn't, because the author got sidetracked working on another project and homework.

Wikipedia. FaceApp — Wikipedia, the free encyclopedia. <http://en.wikipedia.org/w/index.php?title=FaceApp&oldid=1270418493>, 2016.

E. Yudkowsky. Changing emotions. LessWrong, 2009. URL <https://www.lesswrong.com/posts/QZs4vkC7cbyjL9XA9/changing-emotions>.

Ziz. Sincerely, 2017. URL <https://sincerely.fyi>.