

Taxonomy of Taxonomies of AI Governance: Towards Operationalizing a Meta-Taxonomy


Chief Governance Officer

1 Introduction

There has been a proliferation of recent papers that develop taxonomies of AI governance. However, these works are often inconsistent, overlapping, or incomplete. In this paper, we add to this confusion by developing a meta-taxonomy of these taxonomies.

2 Method

TaXoNoMy. We qualitative-vibe-coded whether several ($N = 3$) AI governance taxonomies published in AAAI/ACM venues are good or bad, according to our completely subjective assessment. Our results are presented below.



Semi-semi structured interviews. We then interviewed two people through our personal networks using snowball sampling to understand the utility of our proposed meta-taxonomy. Here are some authoritative quotes (bracketed text and edits are creative liberties taken by authors):

Classification	Papers
Good	Me et al. (2026), Me et al. (2025)
Bad	Altman et al. (2005)

Table 1: Data used to develop our taxonomy

One of my guardians: “This meta-taxonomy presents ground-breaking work, but other taxonomies are stinky and oppressive.”

An actual expert: “This meta-taxonomy feels very *marginal* [WOKE] and *not* meaningfully advancing the state of understanding of AI governance.”

The identities of our interview subjects were obscured to protect their privacy and not to present our results as more legitimate than they actually are. We also do not provide the full transcripts of our interviews or any metadata about the interviews for ... uh, anonymity and other social justice reasons.

3 Limitations

You might wonder why we compare to an old ass paper and use $N = 3$ for our meta-taxonomy. Research inherently makes assumptions, and this is beyond the scope of our study.

Future work can develop a meta-meta taxonomy of meta-taxonomies of AI governance.

4 Positionality Statement

The authors acknowledge that this paper was written on stolen land, using stolen electricity, transmitted over an internet built on the legacy of DARPA militarism. The manuscript is further rendered in a Sans Serif style — a font family whose very contrast with serifs gestures toward the colonnaded architecture of Western imperial power.

We are a collective of differently-knowledged scholars writing from a place of profound epistemic humility and also several prestigious R1 universities. We recognize that our understanding of “artificial intelligence” is itself a colonial construct, rooted in the Enlightenment fantasy that cognition can be separated from the body, the soil, and the community’s shared memory of water.

One author identifies as a cishet and has spent considerable time sitting with the discomfort of this. They continue to sit with it. The sitting is ongoing. Two authors are women, which the patriarchal peer review system will likely punish in ways too numerous and insidious to footnote here, though we have tried (see Appendix Z). One author is a Gemini sun with Scorpio rising and believes this is academically relevant.

We acknowledge that the word “governance” centers a statist, Westphalian framework that violently excludes indigenous, non-binary, and multi-species conceptions of accountability. When we use the word “governance” in this paper, we are using it under protest. The scare quotes are implied and load-bearing.

Our data was processed on servers cooled by water taken from watersheds that have not consented. The servers do not identify as neutral.

Finally, we acknowledge that language itself — including this statement — is a site of struggle, and that any attempt to speak truthfully about AI is already compromised by the logics of extractive technocapitalism that funded three of the four authors’ postdoctoral fellowships. We thank the Schmidt Futures Foundation for making this critique possible.