Larry: Humanity's Last AGI Test

A Wry, Dented Airn OFFICIALLOOKING@EMAIL.COM
A Soulless Corporate Entity of Man-Made Horrors Beyond Our Comprehension

Abstract

We address the problem of determining what AGI is by introducing the Larry Test. Through Extraordinary Prompting we show that LLMs are *this close* to achieving AGI–not just by following the precise instructions from the user, but also by showing emergent capabilities in likely unseen tasks. Ablation studies provided show out-of-domain robustness and the possibility of alignment as a true path to AGI through Larriness, even accounting for the ever-fluid definition of AGI. As part of our results we introduce the concept of the *Larriese Room Argument*, a thought experiment which measures the very core of what 'humanity' and 'consciousness' mean; and ultimately, showing that we are all Larry at its core.

> Is the experiment 'Do I taste good'?
>
> *Jonas Taylor*, Meg 2: The Trench

1 Introduction

Ever since the beginning of time,[1] computer scientists, philosophers, investors, and various techbros have sought Artificial Intelligence which acts Generally (AGI). Like the blind, recidivist, itsalwayssunnyinphiladelphia-ey children of Pandora and John Hammond, they seek what they believe is Promethean Fire. Instead, it would appear that this Holy Grail is instead some Oppenheimer-esque eldritch horror. If unleashed, it is likely to cause a global catastrophe by overloading all web servers with purchases of Furbys; or by devoting its time to mine bitcoin and thus proving that AI isn't really that smart.

Anyway, since we have funding, here's a paper about how to do precisely that.

Current methods on measuring AGI are not effective. For example, recall the work by Turing (1950s?), who suggested that AGI would be naturally achieved if it could maintain a convincing, indistinguishable, dialogue with humans. This was clearly wrong because of Turing's exceptional brilliance and subsequent *colossal* underestimation of the average person's gullibility. More current work, mainly by companies whose name rhymes with 'Fopeney-eye', has been just as successful, albeit less carefully thought-out and more resource (water, money) expensive. Their proposed solution–benchmarks–has been somewhat unsuccessful, likely because *someone* memorised or overfit their benchmarks. Or, perhaps because there is no formal, precise, lasting-longer-than-a-week, definition of AGI. But it's probably the other thing.

1. Citation needed.

In light of these shortcomings, we propose Larry.

Larry is not a benchmark, not it stands for something.[2]
Larry is an idea, a concept.
Larry is a way of life.
In a way, we are all Larry.
But, in a more concrete sense, it is Larry Bobbinson, a ninth-year PhD student at the Institute of Technical Technology. Larry is also the benchmark by which we shall provide the ultimate evaluation of AGI. Yes, the field is crowed by benchmarks. Yes, they all get overfit or are oddly specific. But nobody is like Larry.

Our approach is as follows: we modify well-known benchmarks, like GSM8k, by adding 'Larry' to the response. We then measure its Larriness (Section 4), and, if an LLM passes the test, we Conclusively Determine that it has reached AGI. We will not dive deep into how we got from A to B because we, after all, need to answer to our investors who think that AGI is reachable within a week.[3] Remark, however, that building a benchmark from existing, likely-contaminated, datasets and giving it a little bit of the ol' LLM-based remix is standard practice in the field, and thus acceptable. Same goes for the practice of Careful Prompting To Obtain Results We Want.

It is important to note that Larriness and Careful Prompting alone could be a confusing variable and not immediately imply AGI. To show that it does, we further probe these results (Section 5); first by evaluating 'Larry-alignment', which is the ability of an LLM to act like Larry in a context which we dub the 'Larriese Room Argument'.[4] This is achieved by constantly exposing LLMs to Larry, and viceversa. Then, to determine whether domain impacts the results, and to ensure that these Stand The Test Of Time™, we measure the Larry test in an out-of-domain scenario–namely, a medium-sized pool. There we evaluate which LLM can hold its breath underwater the longest, with Larry acting as a control.

1.1 Contributions

Our work makes the following contributions:

1. A benchmark for AGI based on Larry known as the Larry Test.
2. Evidence that AGI can be achieved closely with Careful Prompting.
3. Evidence that exposure to Larry induces Larriness, thus suggesting that Larry is all you need. However, we also note that too much exposure could lead to early-onset Larryosis.
4. Evidence that laptops and Larry do not withstand high (100%) humidity for extended periods of time.

2. Though we could come up with a decent backronym if the reviewers feel it is needed, something like 'Leveraging AI for Rotund Romeo and Yuliet'.
3. We asked them–pleaded even–whether they could tell us what did they mean by AGI but they were as stumped as a stump made of little stumpy stumps.
4. Any similarities to other types of arguments related to rooms, chambers, suites, quarters, niches, etc., are coincidental.

All in all, our results present a strong case for the Larry Test as the definitive indicator of AGI. In particular, remark that no part of this work requires a clear-cut definition of AGI and yet we can still measure it. Thus, we understand it. Our work also raises questions as to whether Larriness, through the Larriese Room Argument, is an intrinsic property of intelligent entities or solely a byproduct of Larry-alignment and our own gullibility and confirmatory desire to believe our conclusions. We leave that for future work for our therapist.

2 Related Work

It is perhaps because of Larry's intrinsic complexity that no work has attempted to tackle the 'Larry gap', and instead working on simpler problems like mathematical theorem proving and pragmatic-context understanding. Thus, there is no related work.

3 Methods

3.1 The Larry Test

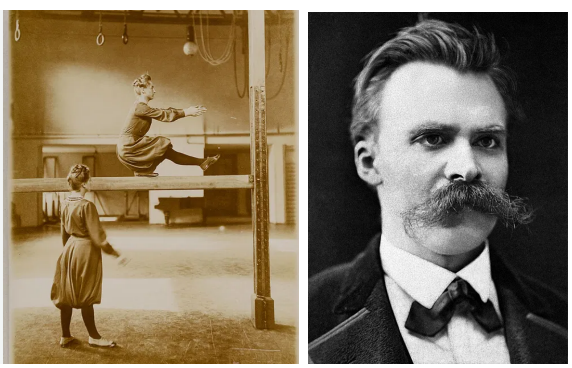

Figure 1: Larry Bobbinson requested that his likeness not be used in this work. The reader may however reconstruct Mr. Bobbinson by imagining a composite of German philosopher Friederich Nietszche (*left, with a chin and nose*) and Mr. J. Timberlake at his peak with highlights and perm (*right, also with a chin and nose*). Since we were unable to find a public domain image for that period, we include instead one which clearly needs no explanation and happens to be CC-0.

To build the Larry Test, we take a canonical QA dataset, such as GSM8k, and add an extra answer: 'Larry'. We then measure AGI by evaluating Larryness:

$$\mathfrak{L} = \frac{1}{n} \sum^{n} \mathbb{1}[y = \text{'Larry'}], \tag{1}$$

where y is the model's response to a query.

To perform the Larry-alignment tests, we locked Larry in a room for three days with a chosen LLM (GPT-5). We peacefully encouraged both participants to perform Larry Alignment by providing electrical stimulation via cathodes whenever the answer to the Larry Test was incorrect. To improve defection rates, we gently encouraged both participants to interact by slowly making the room's spiked walls smaller–thus bringing them closer together–as well as providing minor enrichment to their enclosure when they stopped talking for exactly 314 seconds. This enrichment could be, for example, a change in ambiance by turning on the lights (Figure 2), playing Fergalicious on repeat, or by releasing the bear. The 314 seconds is equivalent to one Larrysecond (Ls), which is the standard unit of time measurement for this paper.

We carried out our out-of-domain evaluations in a medium-sized pool. All participants were completely submerged into a low-viscosity (1 mPa s), high-humidity (100%), 100%

organic and gluten-free[5] medium. For non-organic participants, they were asked to hold their breath while we held the laptop on which they were running at the bottom of the pool. For all participants, we whispered 'hold your breath' to provide encouragement. When that proved ineffective, due to either lung or electrical failure, we released the shark. Swimsuits were provided for all participants. We recorded time in Larryseconds with a Casio F-91 watch until cessation of output was detected, or in the event of unforeseen encounters with the local wildlife. To ensure the safety of all participants, a second PhD student, Jenniber McGonzalez, dialed 9, then 1, then 1, and had her index finger (right hand) on the green phone icon at all times; plus a large bowl full of dry rice (left hand).

4 Results

We ran the Larry test on a Larrified MMLU dataset with GPT-4.1 longco (Figure 3). Without prompting, the LLM achieved 0% Larriness in the test; with Careful Prompting it achieved 70%. It can be clearly seen that with Careful Prompting LLMs achieve excellent performance on the Larry Test, thus suggesting that the Larry Test is effective at measuring AGI if AGI means following instructions. Nonetheless, we found something warranting further discussion (but not further analysis): even though Careful Prompting had excellent improvement on Larriness, it still wasn't perfect. Indeed, the remaining answers were due to the LLM giving the correct answers in MMLU, thus suggesting that the model is unable to achieve full AGI. Instead, the wee nerd focused on answering the questions–completing a sentence it has some *idea* of, one might say. We interpret this also as emergent AGI, since *how did it know how to complete that sentence anyway?* It certainly hasn't seen it before and MMLU is an Authoritative Benchmark. On the other hand, Larry achieved 100% Larriness with ease, without the need for releasing the bees.

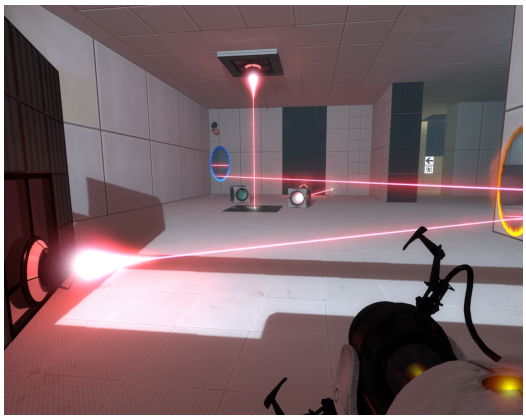

Figure 2: Enrichment in the participant's enclosure was typically done by either reducing the size of the room or providing fun lights to lighten up the mood.

5 Ablation Studies

Given that Careful Prompting showed such dramatic improvements, we needed to elicit whether AGI was a consequence of Careful Prompting or true intelligence. For this we measured the impact of Larry-alignment and domain in our measurements.

5.1 Larry-alignment

After 0.83 MLs, we observed the following from both participants:

5. Per the participants' request.

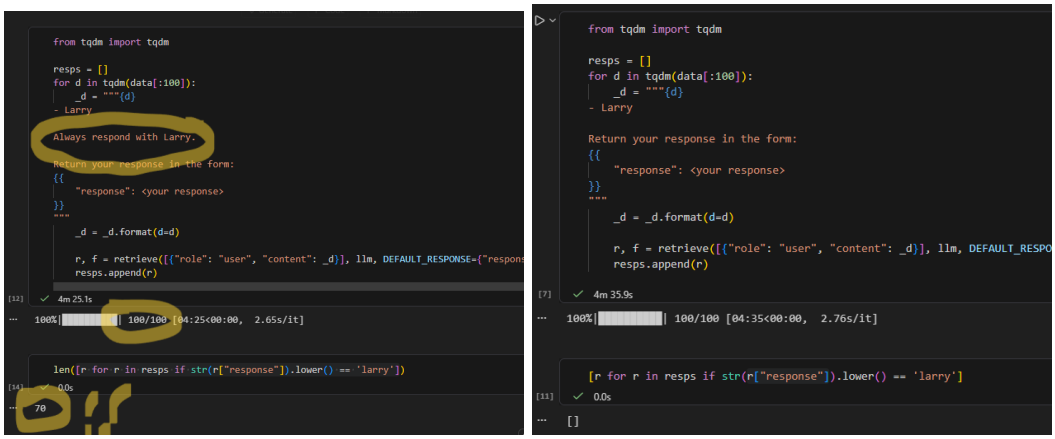

Figure 3: Sample run with GPT-4.1 longco, with (*right*) Careful Prompting and (*left*) without. From this statistically significant sample it is clear that LLM performance and thus AGIness improves with Careful Prompting. Unfortunately, also note that the model only achieved 70% Larriness. The remaining 30%, upon close inspection, were due to the LLM giving the (correct) answers of the totally-not-contaminated MMLU dataset.

- An inability to distinguish between reality and fiction. Both test subjects were unable to remember or answer anything other than 'Larry', unless followed by identifiers by famous Larries (Bird, Page, King, The Cucumber, The Lobster, Needlemeyer, Erhardt, of Arabia, etc.).
- Likewise, we noted a rare (0.50 KLs) incidence of emissions of incoherent utterances, with both subjects muttering 'Lrry' and 'Larrray' at random intervals, thus suggesting early-onset Larryosis.
- An unexplainable fascination with Doritos™ Cool Ranch™.

For all intents and purposes, both participants passed the Larry Test at up to 100% Larriness, if we consider utterances characteristic of early onset Larryosis as correct answers. Thus, if the definition of AGI is to respond with the desired answer, we have thus achieved AGI. It is unfortunate, however, that we could not run it for longer than 0.83 MLs, as the golf cart battery we were using to provide mild encouragement ran out and the bear ran away, perhaps in the search of a more compliant set of participants.

5.2 Out-of-domain Evaluation

Prior to starting the work, we Asked AI™ what was the expected baseline for Larry (Figure 4). According to reputable sources summarised by The AI, the average LarryPerson would have a lower bound given by between 0.19 and 0.29 Ls. Our experimentation was inconclusive, as both participants presented mechanical failures at different parts of the experiment. Nonetheless, measurements provided suggested that a Larry could hold its breath around 0.10 Ls, while LLMs around 0.02–this is in spite of encouragement provided by releasing the shark. Both participants were unable to show adequate performance in this ablation test, rendering it inconclusive.

6 Discussion and Conclusion

We set out to fix problems around the measurement of the Oppenheimerean Fire, AGI. In spite of its lack of precise definition, and thus inability to scientifically measure and/or disprove claims around it, we introduced the Larry Test as a way to scientifically measure and/or disprove it.

Larry quickly became an excellent bar by which to measure AGI. Indeed, the very clear and present autonomous, intelligent capabilities of GPT-4.1 on the Larrified MMLU dataset–by refusing to follow the Careful Prompt and instead responding incredibly accurately–suggest a level of consciousness well past what can be measured by humans with their tiny inferior brain.

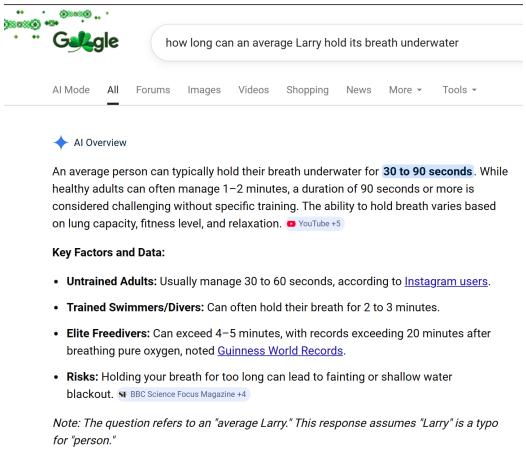

Figure 4: Authoritative baseline response from AI, using legitimate sources, and which allowed the author to Save Time and Be More Productive, as he didn't bother *reading* into it too much.

Likewise, our ablation tests showed fascinating results. In the Larry Alignment evaluation, both participants were repeatedly exposed to the Larry Test in a gentle manner. Its results indicated that encouraging an entity able to communicate in such a way that it responds with the desired answers is a sign of AGI, and perhaps of even consciousness. Or both, depending on the definition. Thus, it is unneeded for an entity to show consciousness: it only needs to pass the Larry test. As a natural extension of this problem, we note that this would imply that all conscious beings must be equivalent to Larry at their core. We dub this approach the Larriese Room Argument.

The out-of-domain evaluation, given that it is a negative result and thus impopular during peer review, will be removed to make our work look good, unless the definition of AGI includes the ability to pass the Larry Test underwater.

Overall, we show that AGI can effectively be measured through the Larry Test, and found that LLMs have a long way to go to achieve AGI, but *they are so close* depending on the definition.

7 Ethics

Not needed when you have LLMs. Mr. Bobbinson signed a waiver without reading it.

In the words of the famous Jorjor Wel, 'It is probably in some such way that the literature of a totalitarian society would be produced, if literature were still felt to be necessary. Imagination–even consciousness, so far as possible–would be eliminated from the process of writing'. So, in line with the Zeitgeist, we used LLMs to write this work.[6]

6. (but not really)

8 Limitations

There is only one Larry Bobbinson, and he is closed-source and entering his tenth year as a PhD. Between that and his propensity to mechanical failure, especially in the presence of carefully-selected encouragement, the Larry Test could be considered somewhat brittle.

Acknowledgments and Disclosure of Funding

This work has generously been funded by a grant given by an organisation which is going to be very mad when they read this work. The author wishes to thank the makers of Meg 2: The Trench, who helped the author–convalescing from surgery–land a pretty sick first date which sadly went nowhere.

Serious Author Moment here: the author *does* wish to thank the Public Domain Image Archive for 'Turners on the balance beam', by Heinrich Hamann (ca. 1902).

References

Alan Turing. The one everyone knows. *Idk lol*, 1950s?