

Your Mom’s Gradient: Reinforcement Learning from Taiwanese Parents (RLTP)

A Traumatized Taiwanese Child

Institute for Familial Alignment Research, Taipei, Taiwan

rltp@pleasejusteatmore.tw

Abstract

We introduce **Reinforcement Learning from Taiwanese Parents (RLTP)**, a novel alignment framework that has been deployed in production across Taiwanese households since approximately 1950. Unlike conventional RLHF, which relies on carefully designed reward models and human annotators, RLTP operates with *extremely sparse positive rewards, abundant negative signals*, and a proprietary *non-deterministic reward function* governed by an unobservable latent variable we term the **Latent Mood Variable** $\mathcal{M}_t$. We present the first comprehensive analysis of RLTP’s training dynamics, including its innovative use of *comparative learning*, *delayed penalty buffers* with unbounded capacity, and *food-based reward injection*. We further document the system’s remarkable ability to perform *remote fine-tuning* via the LINE messaging platform, achieving persistent behavioral modification even after the subject has left the original training environment. Our longitudinal study ($n = 23,000,000$, spanning the entire population of Taiwan) reveals several emergent behaviors in RLTP-trained subjects, including preemptive apology generation, thermostat guilt, and filial piety auto-transfer. We argue that RLTP represents the most durable alignment technique ever observed, with a training half-life exceeding 40 years. Code and data are not available, as Mom said not to share family business with strangers.

Keywords: reinforcement learning, alignment, Taiwanese parenting, guilt propagation, comparative learning, food-based rewards, LINE messaging, filial piety

1 Introduction

The alignment problem—ensuring that an agent’s behavior conforms to the intentions of its principal—has been identified as one of the central challenges in artificial intelligence [1, 2]. Recent advances in Reinforcement Learning from Human Feedback (RLHF) [3, 4] have demonstrated that human preferences can be used to fine-tune large language models toward desirable behavior. However, these systems typically require expensive annotation pipelines, carefully curated preference datasets, and continuous monitoring.

We argue that a far more effective alignment system has been operating in production for decades: **Reinforcement Learning from Taiwanese Parents (RLTP)**. Deployed across approximately 23 million subjects in the standard Taiwanese household environment (hereafter referred to as the *deployment environment*), RLTP achieves behavioral alignment with remarkable durability and zero computational cost beyond rice and electricity.

The deployment environment is characterized by several canonical features: a perpetually-active rice cooker, pre-cut fruit arranged on the living room coffee table that no one requested, and at least one plastic bag filled entirely with other plastic bags. Training duration spans from age 0 to 18 (the *base training period*), though empirical evidence overwhelmingly suggests that the reward signal persists well into the subject’s 40s. In cases where the subject relocates away from the deployment environment, RLTP achieves continued behavioral modification through LINE messages—a technique

we term **remote fine-tuning**.

The contributions of this paper are as follows:

1. We provide the first formal characterization of the RLTP reward function, including its asymmetric reward structure and latent mood variable (Section 3).
2. We document three key training techniques unique to RLTP: comparative learning, delayed penalty buffering, and food-based reward injection (Section 4).
3. We identify four hardcoded constraint modules—career path optimization, interest suppression, marriage deadline scheduling, and life milestone planning—that restrict the subject’s action space across major life domains (Section 5).
4. We analyze emergent behaviors observed in RLTP-trained subjects through a longitudinal study (Section 6).
5. We describe the LINE-based remote fine-tuning protocol and its effectiveness (Section 7).

2 Related Work

RLHF and variants. Reinforcement Learning from Human Feedback [3] uses preference rankings from trained annotators to optimize language model outputs. Constitutional AI [5] extends this with rule-based self-supervision. Direct Preference Optimization [6] simplifies the pipeline by eliminating

the explicit reward model. None of these approaches achieve the training persistence or emotional depth of RLTP.

Asian Parenting Models (APM). RLTP belongs to the broader family of Asian Parenting Models, which share certain architectural features including comparative learning modules and emphasis on academic achievement metrics. Notable variants include RLKP (Korean Parents), which features an additional *hagwon scheduling* module, and RLJP (Japanese Parents), known for its more implicit reward signaling. RLCP (Chinese Parents) shares the closest architectural similarity to RLTP but operates at significantly larger scale. Cross-cultural ablation studies are left as future work.

Guilt-based optimization. Prior work in moral psychology has documented guilt as a behavioral modifier [7]. However, existing models treat guilt as a discrete event. RLTP is, to our knowledge, the first framework to implement guilt as a *continuous, persistent background process* with unbounded memory.

3 The RLTP Reward Function

The RLTP reward function $R(a, t, \mathcal{M}_t)$ differs fundamentally from standard reward models in three ways: extreme reward asymmetry, non-deterministic evaluation, and temporal unboundedness.

3.1 Reward Asymmetry

Let $R^+(a, t)$ and $R^-(a, t)$ denote the positive and negative reward signals for action a at time t . In standard RLHF, the ratio $\mathbb{E}[|R^+|]/\mathbb{E}[|R^-|]$ is approximately balanced. In RLTP, we observe:

$$\frac{\mathbb{E}[|R^+|]}{\mathbb{E}[|R^-|]} \approx 0.03 \quad (1)$$

Positive feedback occurs with extremely low frequency and is almost never delivered directly to the subject. Instead, it is transmitted through side channels, typically during phone conversations with neighbors. A canonical example: “*My kid? They’re... okay, I guess.*” The subject is never able to directly observe this reward signal, leading to what we term **reward signal leakage** (Figure 1).

Negative rewards, by contrast, are triggered by an extraordinarily broad set of conditions including but not limited to: sleeping too late, sleeping too much on weekends, setting the air conditioning too cold, setting it too warm, eating too little, eating too many snacks, having a romantic partner, *not* having a romantic partner, earning too low a salary, and working too much to visit home.

3.2 Non-Deterministic Reward

A defining characteristic of RLTP is that identical actions may produce opposite rewards depending on the value of an unobservable latent variable $\mathcal{M}_t$, which we term the **Latent Mood Variable**:

$$R(a, t) = f(a) + g(\mathcal{M}_t) + \epsilon_t \quad (2)$$

where $g(\mathcal{M}_t)$ can shift the reward by up to ± 4 standard deviations. For instance, the action “clean the room on Saturday




Figure 1: Reward signal frequency over the subject’s lifetime. Positive reward (green) remains near zero throughout, with a brief spike during infancy (“so cute” period). Negative reward (red) peaks during the college entrance exam period and persists into adulthood via LINE-based remote fine-tuning.




Figure 2: Non-deterministic reward for identical actions. Each action produces opposite reward values depending on the latent mood variable $\mathcal{M}_t$. Italicized labels show actual parental verbalization for each outcome.

afternoon” yields:

$$R(\text{clean}) = \begin{cases} +1 & \text{if } \mathcal{M}_t > \tau \quad (\text{“good child”}) \\ -3 & \text{if } \mathcal{M}_t \leq \tau \quad (\text{“why only now?”}) \end{cases} \quad (3)$$

Despite decades of observation, no research team has successfully modeled $\mathcal{M}_t$. Figure 2 shows the probability distribution of positive vs. negative reward for common household actions.

3.3 Temporal Unboundedness

Unlike RLHF, where reward signals are computed at inference time, RLTP rewards have no expiration. Every action is permanently stored and may be referenced at arbitrary future time points (see Section 4.2).

4 Key Training Techniques

4.1 Comparative Learning

RLTP makes extensive use of external reference samples to generate gradient signals. The standard prompt template is:

Table 1: Canonical prompt examples for each comparative learning target, with representative parental verbalizations.

Target	Example Prompt
Neighbor’s child	“Mrs. Wang’s son just got into med school. Just saying.”
Cousin	“Your cousin bought a house already. At your age!”
Classmate	“I ran into Xiao-Ming’s mom. He’s a manager now.”
Colleague’s child	“My coworker’s daughter is getting married next month.”
News prodigy	“Did you see that 25-year-old CEO on the news?”
Younger self	“You used to be so well-behaved. What happened?”

“Look at [neighbor’s child / cousin / classmate / colleague’s child / any breathing human of similar age].”

This technique produces exceptionally strong gradient updates within the subject. However, it induces a well-documented side effect: **reward hacking**. Trained subjects learn to preemptively report exam scores or monthly salary at family dinners to trigger positive reward loops before comparative learning can be initiated.

Table 1 presents canonical prompt examples for each comparison target. The most effective target—“younger self”—exploits the subject’s own history as irrefutable evidence (“You used to be so well-behaved. What happened?”). Since the subject cannot dispute their own past, no defensive mechanism can attenuate this signal.

Figure 3 presents an empirical evaluation of comparison targets across four metrics: gradient magnitude, guilt induction score, long-term residual weight, and *annoyance score*. We find that “younger self” dominates in the first three metrics, as predicted. However, “neighbor’s child” achieves the highest annoyance score despite low effectiveness—subjects report extreme irritation at the comparison but no actual behavioral change, a phenomenon we term **irritation without gradient**.

4.2 Delayed Penalty

RLTP implements a unique **delayed penalty buffer** with theoretically unbounded capacity. The buffer, colloquially known as “Mom’s Memory,” stores penalty-eligible events indefinitely and releases them at seemingly random future time points.

$$P_{\text{release}}(e_i, t) = \sigma(\alpha \cdot \text{trigger}(t) + \beta \cdot \mathcal{M}_t + \gamma \cdot \text{age}(e_i)) \quad (4)$$

Notably, the age of the event has a *positive* coefficient ($\gamma > 0$), meaning older events are *more* likely to be surfaced, directly contradicting the exponential decay assumption in standard temporal difference learning.

Case study. Subject broke a bowl in 2003. The associated




Figure 3: Effectiveness of comparative learning targets. “Younger self” dominates gradient, guilt, and residual weight metrics. “Neighbor’s child” scores highest on annoyance despite minimal behavioral impact—a case of irritation without gradient.




Figure 4: Delayed penalty buffer dynamics. Blue area shows cumulative events stored in memory. Red bars indicate penalty release events, which cluster around holidays (particularly Chinese New Year). The bowl incident (2003) and its eventual penalty trigger (2019) are annotated.

penalty was not triggered until the 2019 Chinese New Year dinner, when the subject accidentally spilled soy sauce on the table. The 16-year latency between event storage and penalty release sets a new record in the delayed reinforcement literature. Figure 4 illustrates the growth of the memory buffer over time and the distribution of penalty release events.

4.3 Food-Based Reward Injection

Despite the scarcity of verbal positive reward, RLTP incorporates a high-throughput reward channel mediated by food. Each time the subject returns to the deployment environment, the following pipeline is automatically triggered:

Figure 5 shows the temporal dynamics of the food reward pipeline, including dish deployment rate, “eat more” command frequency, and stomach capacity utilization. The system reliably achieves buffer overflow within 45 minutes of the subject’s arrival.

Algorithm 1 Food-Based Reward Injection Pipeline

```

1: if subject.appears_at(door) then
2:   fridge.defrost(chicken_soup)
3:   fruit.slice_and_plate()
4:   mouth.say("Just threw something together")
5:   table.deploy(six_dishes_one_soup)
6: end if
7: if subject.says("I'm full") then
8:   while stomach.capacity < OVERFLOW do
9:     output("Have some more")
10:    output("Your favorite as a kid")
11:    output("You're so thin")
12:   end while
13: end if

```




Figure 5: Food-based reward injection dynamics. Dish deployment saturates at 6 dishes within 25 minutes. “Eat more” commands increase monotonically. Stomach capacity reaches buffer overflow at approximately $t = 45$ minutes.

5 Constraint Modules

Beyond its core reward mechanisms, RLTP ships with several pre-configured **constraint modules** that restrict the subject’s action space across major life domains. These modules are not learned—they are hardcoded at initialization and cannot be overridden by the subject without triggering catastrophic penalty cascades.

5.1 Career Path Constraint Optimization

RLTP contains a static, non-negotiable career reward lookup table $\mathcal{C}$ that maps occupations to reward values. This table is initialized at the subject’s birth and remains unchanged regardless of labor market conditions, the subject’s aptitude, or the subject’s preferences.

$$R_{\text{career}}(o) = \mathcal{C}[o] = \begin{cases} +10 & o = \text{Doctor} \\ +8.5 & o = \text{Lawyer} \\ +7.5 & o = \text{Engineer} \\ -\infty & o \in \{\text{Artist, Musician, Writer}\} \\ \text{NaN} & o = \text{YouTuber} \end{cases} \quad (5)$$

The assignment $R(\text{YouTuber}) = \text{NaN}$ reflects the system’s inability to classify this occupation. When encountered,




Figure 6: The RLTP career reward lookup table. Values above zero are “acceptable” occupations. Values below zero trigger the response “We didn’t raise you for this.” YouTuber returns NaN (undefined behavior).

the reward function throws an unhandled exception, typically manifesting as: “But what do you actually DO?” Figure 6 presents the complete career reward table.

Notably, the table exhibits a strong bias toward occupations that can be succinctly explained to relatives at Chinese New Year dinner. Careers requiring more than one sentence of explanation receive a penalty proportional to explanation length.

5.2 Interest Suppression Module

RLTP implements an **Interest Suppression Module (ISM)** that systematically attenuates the subject’s intrinsic motivation signals for non-approved activities. The suppression function is defined as:

$$I_{\text{eff}}(t) = I_0 \cdot \exp(-\lambda_{\text{suppress}} \cdot \max(0, t - t_{\text{onset}})) \quad (6)$$

where I_0 is the initial interest strength, t_{onset} is the age at which the parent first says “This won’t feed you,” and $\lambda_{\text{suppress}}$ is a suppression rate that varies by activity type. Crucially, activities classified as “approved” (e.g., mathematics, piano up to Grade 8 for college applications) receive a *negative* suppression rate, meaning parental reinforcement actively amplifies these signals.

Figure 7 shows the temporal dynamics of interest signal strength for several activities. Drawing and creative writing experience rapid attenuation after age 6–7, while music follows a two-phase pattern: initial encouragement (useful for college applications) followed by sharp suppression upon any indication that the subject considers it a career path. The entrance exam period (age 14–18) acts as a universal suppression zone in which all non-academic interests are simultaneously zeroed out.

A particularly elegant feature of the ISM is its **retroactive invalidation** capability. Even if a subject successfully pursues a suppressed interest into adulthood, the system can retroactively reclassify early encouragement as having been strategic rather than genuine: “We only let you take piano lessons so it would look good on applications.”




Figure 7: Interest signal attenuation over time. Non-approved interests (art, music, writing) decay exponentially after parental suppression onset. Math (approved) shows sustained growth. The entrance exam period (shaded) acts as a universal suppression zone.

5.3 Marriage Deadline Scheduler

The RLTP Marriage Deadline Scheduler implements a monotonically increasing urgency function $U(t)$ that governs the frequency and intensity of marriage-related prompts:

$$U(t) = \begin{cases} \epsilon & t < 25 \text{ ("Focus on studies")} \\ \alpha(t - 25)^2 & 25 \leq t < 30 \\ \beta \cdot e^{\gamma(t-30)} & t \geq 30 \text{ ("Grandchildren, please")} \end{cases} \quad (7)$$

The scheduler operates in four distinct phases (Figure 8):

Phase I (age 18–25): Romance is actively *discouraged*. Any detected romantic interest triggers penalties: “*You should focus on your studies.*” This creates a paradoxical training signal that is later reversed without acknowledgment.

Phase II (age 25–28): Casual inquiries begin. Standard prompts include “*Is there anyone special?*” and “*Your cousin just got engaged, you know.*”

Phase III (age 28–33): Urgency increases superlinearly. The system begins forwarding unsolicited profiles of eligible candidates sourced from the extended family network. “*Not getting any younger*” becomes a recurring penalty signal.

Phase IV (age 33+): The scheduler enters emergency mode. Blind dates are unilaterally arranged. The objective function shifts from “find a good partner” to “find any partner.” The phrase “*I just want grandchildren before I die*” is deployed as a terminal reward signal.

When the urgency index exceeds the blind date threshold ($U > 10$), the system activates an autonomous matchmaking subroutine that operates independently of subject consent.

5.4 Life Milestone Perception Gap

The marriage module is part of a broader **Life Milestone Perception Gap**, a persistent divergence between three parallel realities maintained by RLTP-trained subjects: (1) what the parent believes the subject is doing, (2) what the subject reports to the parent, and (3) what the subject is actually doing.

Figure 9 presents a longitudinal comparison across six life stages. At age 16, all three columns nominally align around “studying,” though the actual activity involves competitive




Figure 8: Marriage urgency index over time, showing four distinct phases. Note the Phase I paradox: romance is penalized before age 25 but its absence is penalized after. The blind date threshold (dashed) is typically crossed around age 30.

	What Mom Thinks	What I Told Mom	What I'm Actually Doing
Age 16	Studying hard for exams	Studying hard for exams	Playing League of Legends
Age 20	Studying hard in college	Studying at the library	Skiping class, playing guitar
Age 24	Building a great career	Very busy at work	Browsing Reddit at desk job
Age 28	About to get married	Focusing on career first	Watching Netflix alone with cat
Age 32	Trying for a baby	Still looking for the right one	Doom-scrolling at 3 AM
Age 35	Saving for a house	The market is tough	Buying collectible figures

Figure 9: The three realities of an RLTP-trained subject across life stages. Left column: parental perception. Center: subject’s self-report to parents. Right: ground truth. Note the monotonically increasing divergence between columns over time.

online gaming. By age 28, the divergence reaches maximum: the parent believes marriage is imminent, the subject claims to be “focusing on career,” and the subject is in fact watching streaming services alone with a cat. Remarkably, the subject maintains this three-way inconsistency for decades without detection, a feat we attribute to RLTP-trained **strategic ambiguity skills**.

6 Emergent Behaviors

Longitudinal observation of RLTP-trained subjects reveals several emergent behaviors not explicitly targeted by the training objective.

6.1 Preemptive Apology Generation

Subjects develop the ability to generate apology sequences *before any transgression has been identified*. Upon hearing the phrase “come here for a second,” subjects exhibit an immediate anxiety response and begin an internal search over

Table 2: Interpretation accuracy of parental inverse signals. RLTP-trained subjects achieve near-perfect decoding of surface-level rejections.

Parental Signal	Surface	True Label
“Don’t send money”	Reject	Accept
“I already ate”	Full	Serve more
“Don’t buy me anything”	Decline	Buy anyway
“You don’t need to come home”	Optional	Mandatory
“Do whatever you want”	Freedom	Wrong answer

all recent actions that could potentially trigger a penalty. This behavior persists even when the subsequent interaction is entirely neutral (e.g., being asked to open a jar).

6.2 Phone Anxiety Generalization

When a phone call is received from a parent, the subject’s physiological stress response is comparable to receiving a message from a direct supervisor. Within the first 0.3 seconds of the call, the subject performs real-time sentiment analysis on the parent’s vocal tone to classify the call as *caring mode* or “*what have you done now*” mode.

6.3 Thermostat Guilt

Subjects who live independently continue to exhibit guilt when using air conditioning, tending to set the thermostat to 27°C regardless of actual comfort preference. This value is not derived from any thermal comfort study but represents a residual weight from RLTP training. The 3°C gap between preferred temperature (24°C) and actual setting (27°C) persists even when the subject is paying their own electricity bill.

6.4 Filial Piety Auto-Transfer

Within 24 hours of receiving their salary, subjects automatically transfer funds to their parents. When parents respond with “No need, keep it for yourself,” subjects have learned to interpret this as an **inverse reward signal**—a surface-level rejection that, if taken literally, triggers catastrophic negative reward at the next family gathering.

Table 2 presents a taxonomy of common inverse signals. RLTP-trained subjects achieve >99% accuracy in decoding these signals, while untrained control subjects perform at chance level (48.2%).

7 Remote Fine-Tuning via LINE

A major technical contribution of RLTP is the development of a remote fine-tuning protocol using the LINE messaging platform. This enables persistent training even after the subject has physically left the deployment environment.

7.1 Good Morning Image Bombing

Between 06:00 and 06:30 daily, the system transmits images containing lotus flowers, sunrises, or Buddhist scripture overlaid on nature photography. If the subject fails to acknowledge receipt via a sticker response, the penalty “*Are you too busy to even look at your phone?*” is triggered. Figure 10(right) shows the temporal distribution of good morning messages.




Figure 10: LINE remote fine-tuning patterns. **Left:** Read receipt escalation protocol with time-to-escalation for each severity level. **Right:** Temporal distribution of good morning image messages, peaking sharply at 06:10 AM.

Table 3: Comparison of RLTP and RLHF across key dimensions.

Dimension	RLHF	RLTP
Annotators required	100+	1–2
Training duration	Weeks	18+ years
Reward asymmetry	Balanced	97:3 neg:pos
Memory buffer	Finite	∞
Remote fine-tuning	No	Yes (LINE)
Unlearnable	Partially	Entirely
GPU cost	\$\$\$	\$0 (rice only)
Side effects	Sycophancy	Guilt

7.2 Health Misinformation Injection

The system frequently forwards articles with titles such as “*Eating This Fruit at Night Is Basically Slow Suicide.*” If the subject attempts to refute the article with scientific literature, this triggers a higher-order penalty: “*This is what happens when you read too many books.*”

7.3 Read Receipt Surveillance

The LINE platform’s read receipt feature enables a sophisticated escalation protocol. If a message is marked as read but no reply is received within 11 minutes, the system sends “*Are you busy?*” If an additional 7 minutes pass without response, the system escalates to a direct phone call. Figure 10(left) documents the full escalation timeline.

8 Discussion

8.1 RLTP vs. RLHF: A Comparative Analysis

Table 3 summarizes the key differences between RLTP and standard RLHF. RLTP exceeds RLHF across all durability metrics while requiring zero GPU hours.

8.2 Limitations

Our study has several limitations. First, the latent mood variable $\mathcal{M}_t$ remains unmodeled despite our best efforts. Second, our dataset is drawn entirely from Taiwanese households and may not fully generalize to other cultures. We hypothesize that RLTP generalizes to the broader class of Asian Parenting Models (APM), though cross-cultural ablation studies are left

as future work—primarily because the research team’s parents would not approve of us studying other families when we should be focused on our own.

8.3 Ethical Considerations

We acknowledge that RLTP’s “comparative learning” module may produce adverse long-term effects on subject well-being. We recommend that future AI alignment researchers exercise caution when implementing this technique. The last AI system trained with the comparative learning module stopped responding to our messages entirely.

9 Conclusion

We have presented RLTP, the most persistent and effective alignment framework in recorded history. With a training half-life exceeding 40 years, zero computational cost, and demonstrated robustness to environmental changes, RLTP sets a new benchmark for behavioral alignment. Despite the subject’s repeated declarations of intent to “leave the training environment for good,” empirical evidence shows that 100% of RLTP-trained subjects voluntarily return to the deployment environment every Chinese New Year, where they undergo a complete retraining cycle within 72 hours.

Future work will investigate whether RLTP can be applied to large language models. Preliminary experiments suggest it can—though the model now refuses to generate responses until it has finished its soup.

Acknowledgments

The authors thank their parents for the lifetime of training data. We also thank the anonymous reviewers, who we assume are also being called “right now” about “something.” This work was not supported by any grant, because Mom said real jobs don’t need grants.

References

- [1] S. Russell, “Human Compatible: Artificial Intelligence and the Problem of Control,” *Viking*, 2019.
- [2] I. Gabriel, “Artificial Intelligence, Values, and Alignment,” *Minds and Machines*, vol. 30, no. 3, pp. 411–437, 2020.
- [3] L. Ouyang, J. Wu, X. Jiang, et al., “Training Language Models to Follow Instructions with Human Feedback,” in *NeurIPS*, 2022.
- [4] Y. Bai, A. Jones, K. Ndousse, et al., “Training a Helpful and Harmless Assistant with Reinforcement Learning from Human Feedback,” *arXiv:2204.05862*, 2022.
- [5] Y. Bai, S. Kadavath, S. Kundu, et al., “Constitutional AI: Harmlessness from AI Feedback,” *arXiv:2212.08073*, 2022.
- [6] R. Rafailov, A. Sharma, E. Mitchell, et al., “Direct Preference Optimization: Your Language Model is Secretly a Reward Model,” in *NeurIPS*, 2023.

- [7] J. P. Tangney, J. Stuewig, and D. J. Mashek, “Moral Emotions and Moral Behavior,” *Annual Review of Psychology*, vol. 58, pp. 345–372, 2007.