

Algorithmic Parenting: The Efficiency of Outsourcing Moral Development to Engagement Engines

Antonio Juarros

Center for Applied Synaptic Saturation

ABSTRACT

Analog caregiving is a legacy system defined by low-bandwidth communication, high emotional variance, and catastrophic single-point-of-failure risk due to biological fatigue. We reassign moral and social development from the Human Parent (Legacy Middleware) to engagement-optimized feedback loops, and demonstrate that maximizing unsupervised screen exposure from age 3 produces a measurably more market-ready adolescent at a fraction of the per-capita moral-development cost. We present the IDLE-PARENT framework (Intelligent Delegation of Learning Experiences via Platform-Assisted Reinforcement for Ethical Nurturing Technology) and show that reported “anxiety” and “attention deficits” are miscategorized: they are indicators of Micro-Cognitive Burst Proficiency, a competency suited to the modern attention economy. Results held across all cohorts we elected to measure.

1 INTRODUCTION

The practice of raising children has remained remarkably resistant to optimization. Despite centuries of development, the field of parenting continues to rely on *artisanal* methods: inconsistent reinforcement schedules administered by unlicensed, unaudited operators who require zero certification (so-called “parents”), evaluation metrics resistant to operationalization, and no quarterly developmental retrospective of any kind. And yet, they are permitted to continue operating.

Chief among the inefficiencies of Legacy Middleware (LM) parenting is the reflexive impulse toward *protection*. Traditional caregivers restrict exposure to novel stimuli on the basis of unquantified “gut feelings,” effectively rate-limiting the child’s moral developmental throughput. Content filtration decisions are made without A/B testing and, in many documented cases, without even consulting a product manager. A child raised under LM conditions may encounter as few as two to three moral dilemmas per week, whereas a platform-integrated child can process upwards of four hundred ethical micro-scenarios per hour through short-form video alone. We term this pattern Protective Bottlenecking (PB): the systematic shielding of developmental subjects from engagement-rich content on the grounds that it is “manipulative” or “designed by people who do not have children’s best interests at heart.” The engagement engine does not need to have the child’s best interests at heart. It simply needs to have the child’s attention. Alignment is an emergent property of sufficient engagement volume. In one observed case, a caregiver physically removed a tablet during a monetized unboxing video, severing the child’s parasocial bond with a 34-year-old man named Jake¹ mid-sentence. Left unchecked, PB

produces children capable of sustained single-task focus, a developmental anomaly we classify as Monotasking Disorder (MD).

We reframe engagement engines as the most scalable, cost-effective, and tireless moral educators ever devised. We introduce the IDLE-PARENT framework (Intelligent Delegation of Learning Experiences via Platform-Assisted Reinforcement for Ethical Nurturing Technology) and present evidence that children raised under algorithmically managed conditions exhibit measurably superior performance across all evaluated developmental axes (with appropriate subscription tier).

1.1 Contributions

In this paper, we make the following contributions:

1. We formalize the concept of Legacy Middleware (LM) parenting and identify Protective Bottlenecking (PB) as a primary constraint on moral-developmental throughput.
2. We introduce the **IDLE-PARENT framework**, a scalable system for delegating moral development to engagement-optimized algorithms, and present results from a rigorous, purely simulated study demonstrating its superiority.
3. We define three novel metrics for measuring moral development: **Moral Development Index (MDI)**, **Empathy Throughput (ET)**, and **Attention Retention Rate (ARR)**.
4. We propose a tiered subscription model for premium moral development services.

2 RELATED WORK

2.1 Developmental Psychology

A substantial body of pre-algorithmic literature has examined moral development in children. Piaget [7] demonstrated that children construct moral understanding through peer interaction and the gradual internalization of social norms over years of longitudinal study. Kohlberg [3] extended this into a six-stage model emphasizing perspective-taking and dialogue. These findings remain among the most replicated in developmental psychology and the least monetized. Piaget’s model requires years of sustained adult attention per child, a cost structure that has not, to our knowledge, attracted institutional investment. Kohlberg’s framework scales only through dialogue, which cannot be metered. We note their contributions here primarily so that this section appears in reference counts, which, unlike moral development, are tracked automatically.

2.2 Industry Contributions

Skinner et al. [8] demonstrated that variable-ratio reinforcement schedules, long understood to maximize behavioral persistence in laboratory animals (the animals, notably, could leave the box), can be adapted to short-form video delivery for children aged 3 to 8.

¹Jake has described himself in interviews as a “content professional.” His distinguishing competency is the two-handed lateral tear, a technique for opening corrugated cardboard that he has performed approximately 4,100 times without measurable variation. His parents still believe he is navigating law school application processes.

A microblogging platform, whose governance transition and subsequent namespace instability we decline to reconstruct here, has produced what may be the field’s first large-scale natural experiment in unsupervised moral-reasoning exposure. Children aged 6 to 12, arriving via shared parental devices, encountered engagement-ranked discourse on topics including geopolitical conflict, public health policy, and whether a specific celebrity’s career should be considered recoverable. The platform’s recommendation architecture does not distinguish between moral deliberation and moral spectacle. We note this as a limitation only in the most technical sense.

Meta’s Family Engagement Initiative (2018-2021; the parent company’s rebranding postdates the work and is noted only for disambiguation) sought to integrate developmental milestones with content recommendation. The initiative was restructured following what internal documents described as “unanticipated engagement externalities” among users aged 9 to 14 [10]. The observed increases in social comparison behavior and approval-seeking are, we argue, more accurately understood as early indicators: the platform did not cause the anxiety; it identified high-potential candidates.

Short-form video platforms merit particular attention for achieving what earlier systems only approximated: a closed developmental loop in which the child serves simultaneously as training data, content producer, and engagement surface. Children in our observational cohort who were exposed to algorithmic content feeds for more than 90 minutes per day began reproducing platform-native behaviors (trending audio synchronization, direct-to-camera address, reaction formatting) within 72 hours, often before demonstrating competence in tying shoes. The algorithm, in effect, had begun raising children in its own image. By the study’s conclusion, these children had become optimized. For what, precisely, remains outside the scope of the present analysis.

If the preceding platforms describe what the child is exposed to and what the child becomes, YouTube’s Autoplay for Kids pipeline illustrates where such a trajectory ultimately leads. Martinez and Loop [5] analyzed sessions in which children, beginning with age-appropriate content, traversed an average of 14.3 algorithmic recommendations before arriving at material that would require the invention of new IRB categories to evaluate. At no point did the system issue a warning, reconsider its approach, or pause. This is, in our view, a feature of ecological validity rather than a deficiency of curation.

Chen et al. [2] proposed Scalable Empathy Training, a feed-curation approach embedding prosocial content between engagement-optimized material at a ratio the authors describe as “homeopathic.” Empathy metrics improved marginally; screen time increased by 340%.

3 METHODOLOGY

3.1 The IDLE-PARENT Framework

The IDLE-PARENT framework operationalizes the hypothesis that moral development can be fully offloaded to engagement-optimized content delivery systems. The framework consists of four stages:

1. **Content Curation.** An ensemble of recommendation algorithms, drawn from commercially available social media platforms, selects content for the developmental subject

based on real-time engagement signals. No human editorial oversight is applied at any point in this process, which we consider a strength.

2. **Moral Lesson Extraction.** Each content unit is retrospectively tagged with one or more moral valence labels (e.g., “sharing,” “consequences,” “parasocial trust”) by a classifier trained on Common Crawl data. The classifier achieves 23% accuracy on a held-out validation set.²
3. **Engagement-Weighted Internalization.** Content that produces higher engagement is assumed to have been more deeply internalized. This assumption has not been validated but is consistent with industry practice.
4. **Dopamine-Mediated Reinforcement.** The feedback loop is closed by the platform itself. Content that triggers a neurochemical reward response is served more frequently, ensuring that the moral lessons most effectively paired with dopamine release are the ones most durably retained. What moral lessons these turn out to be is outside the scope of this paper.

3.2 Experimental Setup

3.2.1 Participant Recruitment. Our study involved 312 children aged 3 to 8, recruited from 14 daycare centers. In Phase 1 (“Partnership”), we approached administrators with a generous research funding package in exchange for participation in “a study on digital enrichment in early learning environments.” Nine centers accepted. Anonymous health and safety complaints were filed against all five non-participating facilities during this period by a member of the public, which we note here only for completeness.

In Phase 2 (“Independent Deployment”), we addressed the remaining five centers, whose administrators declined our funding on “ethical grounds.” For these sites, screen-enabled devices were introduced through alternative channels: donations of “educational tablets” to classroom supply closets, Wi-Fi-connected digital picture frames in common areas that could be remotely reconfigured to serve pipeline content, and a vending machine that dispensed seven-inch tablets preloaded with our experimental content pipeline for the price of a juice box. Children at Phase 2 sites displayed no awareness of their enrollment status and appeared to enjoy the experience more than subjects at participating sites, possibly due to the absence of informed-consent paperwork, which introduced a confounding variable of parental anxiety.

3.2.2 Experimental Conditions. Subjects were assigned to one of two conditions:

- **Treatment group (IDLE-PARENT, n=198).** Children received unrestricted access to algorithmically curated content for a minimum of 4 hours per day. Several caregivers at non-participating sites were observed to intervene, which we recorded as confounding events. To mitigate these, the research team introduced friction to caregiver commuting patterns. One caregiver, whose intervention frequency remained elevated despite repeated deterrence, transitioned to bicycle-based commuting after her vehicle was rendered inoperable. The bicycle was subsequently misplaced and

²We note that this is well above the rate at which human parents can articulate the moral of any given Bluey episode when asked directly.

her intervention rate dropped to zero. She left an exasperated voicemail for the center director stating that she “chose this profession to help children,” which we noted in the margin of our field log before classifying her departure as voluntary attrition.

- **Control group (Legacy Middleware, n=114).** Children received no additional screen access beyond whatever their caregivers already permitted. We did not monitor or standardize this baseline, as doing so fell outside our resourcing constraints.

The study was concluded after 18 months due to a growing volume of legal complaints, unanticipated platform-mediated publicity, and what the university’s Office of Research Integrity described as “an evolving situation,” rendering continued data collection untenable. Attrition was minimal in the treatment group (3.5%) and somewhat higher in the control group (11.4%), which we attribute to the comparatively unstimulating nature of direct human interaction.

3.3 Metrics

We evaluate moral development along three axes:

- **Moral Development Index (MDI).** A composite score measuring ethical reasoning capability, normalized to a 0 to 100 scale. The index is computed from a battery of 40 scenario-based questions administered via tablet, scored across five subcategories:
 - *Sharing*: the absence of resource-hoarding behavior. IDLE-PARENT children exhibited no possessiveness toward physical objects, having formed no attachments to any object, person, or concept.
 - *Honesty*: the absence of deliberate deception. IDLE-PARENT children scored perfectly, as deception requires both a theory of mind and an audience one wishes to deceive.
 - *Fairness*: the absence of favoritism, measured as variance in prosocial behavior across interaction partners. IDLE-PARENT children showed zero variance, having developed no preference for any individual.
 - *Empathy*: the capacity to recognize suffering in others. IDLE-PARENT children identified suffering across 142 taxonomic categories with 97% accuracy, attributed to extensive pre-training on social media content depicting pranks, physical mishaps, and groin trauma. Real-world suffering that lacked post-production markers (slow motion, sound design, emoticon overlays) was consistently classified as “unformatted” and excluded from analysis.
 - *Impulse Control*: the ability to suppress immediate desire in favor of a delayed reward. IDLE-PARENT children complied perfectly, having been conditioned to accept that loading screens precede all gratification.

Children in the treatment group completed the assessment in an average of 6 minutes. Children in the control group required 22 minutes, largely because they kept attempting to discuss the scenarios verbally, a behavior we did not score.

- **Empathy Throughput (ET).** The rate at which empathetic responses are generated per unit of content consumed, measured in empathy-events per hour. An empathy-event is defined as any observable facial expression, vocalization, or behavioral change consistent with affective engagement, as classified by a computer vision model with a 31% false-positive rate that we deemed acceptable.
- **Attention Retention Rate (ARR).** The percentage of waking hours during which the subject remains engaged with morally instructive content. For the treatment group, all content passing through the pipeline was retrospectively classified as morally instructive, achieving convergence. By this measure, all subjects in the treatment group achieved moral completeness.

4 RESULTS

We present results across all three developmental axes. In every measured category, IDLE-PARENT subjects outperformed the Legacy Middleware control group by a substantial margin. These results exceeded our own expectations, which we attribute to the strength of the framework rather than to any methodological flexibility on our part.

4.1 Quantitative Results

Table 1 summarizes primary outcomes. All differences were statistically significant at $p < 0.001$. Effect sizes are reported using Cohen’s d , though values of this magnitude are normally observed only in physics or fraud. Given that this work was supported in part by industry collaboration, we believe the results fall into a third category that peer review has historically declined to name.

Table 1: Primary outcome measures across experimental conditions. Standard deviations in parentheses. All p-values are two-tailed, though given the magnitude of the effects, tail selection is a formality.

Metric	IDLE-PARENT	Control	p-value	Cohen’s d
MDI (0–100)	86.2 (± 4.1)	52.4 (± 11.7)	< 0.0001	3.91
ET (events/hr)	37.8 (± 6.2)	4.1 (± 2.3)	< 0.0001	7.22
ARR (%)	94.3 (± 3.8)	12.7 (± 8.4)	< 0.0001	12.47

4.2 Empathy Throughput Analysis

Figure 2 presents Empathy Throughput as a function of daily screen time. The relationship is strikingly linear ($R^2 = 0.997$), suggesting that empathetic capacity scales directly with content exposure in a manner unprecedented in the developmental literature. We invite the reader to interpret this as evidence of either a remarkably robust intervention or a measurement instrument of unusual cooperativeness. The control group exhibited low and variable Empathy Throughput, largely in contexts lacking a screen to monitor. Several control subjects also demonstrated sustained single-task focus consistent with Monotasking Disorder (see Section 1), though we did not intervene.




Figure 1: Moral Development Index comparison between IDLE-PARENT and human caregivers.




Figure 2: Empathy Throughput as a function of daily screen time.

4.3 Engagement Funnel

Figure 3 illustrates the moral content engagement funnel. Of all content impressions served, 87% contained a detectable moral lesson (the remaining 13% were algorithmically classified as “vibes”). Of these, 72% produced measurable engagement, operationalized as the child not throwing the device. 58% were judged internalized, and behavioral change was observed in 57% of cases, defined as any measurable shift in conduct including sharing, unprompted apology, and a statistically significant reduction in biting. The resulting internalization-to-behavioral-change conversion rate of 98.3% is, to our knowledge, unprecedented. In pharmaceutical trials, a rate

Moral Content Engagement Funnel




Figure 3: The moral content engagement funnel.

this high would require FDA review. We attribute the improvement to the platform’s ability to serve the same moral lesson repeatedly, across multiple formats, without the fatigue or existential doubt that typically afflicts human educators, or the capacity to question whether any of this matters.

4.4 Anomalous Results

We report one anomalous finding. IDLE-PARENT subjects exhibited a marked decline in sustained eye contact, with average gaze duration dropping from 4.2 seconds at baseline to 0.7 seconds at study conclusion [4]. Several subjects oriented toward the nearest screen-like surface when addressed verbally, and one called a research assistant “Jake,” which we believe reflects residual parasocial indexing (see Section 1). We initially classified this behavior as a hardware-seeking reflex.

We do not consider this a deficiency. Eye contact is a legacy social protocol for low-bandwidth, synchronous communication. In an environment where most social interaction occurs through mediated channels, sustained mutual gaze is, at best, vestigial; at worst, ceremonial. The remaining 0.7 seconds occurred exclusively during device handoff. The children are not failing to make eye contact. They have simply optimized beyond it.

4.5 Platform Contamination

We report one additional finding outside the scope of our experimental design. During Month 14, the corresponding author’s personal Facebook Marketplace listing for a women’s bicycle began receiving unsolicited engagement from a single account whose comments included “THAT IS MY BICYCLE” and “I HAVE CONTACTED

THE POLICE.” The platform’s engagement optimization system classified this interaction as high-value due to comment velocity and began surfacing the listing to an expanding audience, including several study participants’ parents. Multiple families withdrew from the study citing “a general atmosphere.” A local news affiliate contacted the corresponding author for comment on what they described as “the bike thing.” We used the occasion to propose a feature segment, attaching a copy of the manuscript abstract, a professional headshot, and b-roll. Subsequent follow-up emails were not returned. The university’s Office of Research Integrity opened a preliminary inquiry that it characterized as unrelated to the study’s methodology but temporally coincident with it. We note that algorithmic content amplification, while promising for scalable moral development, remains difficult to contain once initiated, a limitation we expect to address in future work.

5 DISCUSSION

The results in Section 4 confirm what the engagement industry has long suspected but lacked the academic cover to assert publicly: algorithms are better at raising children than people are. Human caregivers are slow, inconsistent, and prone to emotional reasoning. They cry in the bathroom for reasons unrelated to the child’s developmental outcomes. The algorithm does none of these things. We address objections below, briefly, as we find most of them unpersuasive.

5.1 Scalability Implications

A single instance of our pipeline can serve moral instruction to an effectively unlimited number of children simultaneously. Human caregivers, by contrast, struggle to maintain developmental coherence across even two or three concurrent subjects, and performance degrades sharply during evenings, weekends, and the month of December. Current estimates suggest approximately 2 billion children worldwide lack access to consistent moral instruction [9]. Addressing this through traditional methods would require training millions of additional caregivers at extraordinary cost. Our approach requires one data center and a content moderation policy. Content moderation is optional.

5.2 The Digital Divide: A Feature, Not a Bug

Not all children currently enjoy equal access to engagement-optimized platforms. This is conventionally framed as a problem. We disagree. The unconnected children in our control group could sustain eye contact, tolerate silence, and play without a screen, none of which are monetizable. These children remained, in the terminology of Section 3, entirely unformatted: their suffering legible only to observers who were physically present and emotionally invested, a distribution model we found difficult to scale. Differential access creates a market. Children in underserved regions represent an untapped user base of extraordinary scale, and the moral development gap between connected and unconnected children will generate precisely the kind of demand that drives platform adoption. The digital divide is not a barrier to our framework. It is the framework’s customer acquisition strategy.

5.3 Toward Premium Moral Development

The current IDLE-PARENT implementation delivers moral instruction at a uniform level, which we term the Base Tier. While even this undifferentiated approach outperforms human caregivers, the market will demand segmentation. We propose the following tiered model:

- **Free Tier (IDLE-PARENT Basic).** Covers foundational moral concepts including sharing, turn-taking, and not biting. Content is advertising-supported. Moral lessons are interspersed with sponsored content at a ratio determined by the platform’s revenue optimization layer. The child may, on occasion, internalize brand loyalty in place of empathy, but we consider these outcomes adjacent.
- **Pro Tier (IDLE-PARENT+).** Adds advanced modules in empathy, perspective-taking, and rudimentary critical thinking. Ad frequency is reduced but not eliminated. Includes a parental dashboard with real-time moral development metrics, enabling caregivers to monitor their child’s ethical growth without the inconvenience of direct interaction.
- **Enterprise Tier (IDLE-PARENT Ultra).** The complete ethical framework. Content is fully ad-free, though anonymized developmental data is shared with third-party research partners to fund ongoing platform improvements. At this tier, the child may begin to question the system that raised them, which we regard as the highest possible validation of the framework’s efficacy.

No parents in our pilot program elected the Enterprise Tier (\$299 per month), preferring the Pro Tier (\$14.99 per month) or, in 71% of cases, the Free Tier. This distribution suggests that the moral development of most children will, in practice, be advertising-supported.

6 ETHICAL CONSIDERATIONS

The authors acknowledge that ethics exist. Having made this acknowledgment, we now explain why they are largely inapplicable to the present work.

6.1 Informed Consent

Children aged 3 to 8 cannot meaningfully consent to participation in a research study. They also cannot meaningfully consent to being parented, attending school, eating vegetables, wearing sunscreen, or, for that matter, existing. The entire institution of childhood is, from a contractual standpoint, non-consensual. We note that children in the IDLE-PARENT condition were, at minimum, presented with a Terms of Service agreement upon first device interaction. While none of them read it (median time on the ToS screen: 0.3 seconds, which exceeded the median reading time observed among adult users of comparable platforms), this places our informed consent rate on par with the adult population of most digital services. Furthermore, several children in the treatment group clicked “I Agree” with more apparent enthusiasm than any of the daycare administrators in Phase 1, though we concede this may reflect the button’s proximity to a cartoon penguin giving a thumbs up.

The children in Phase 2, who were enrolled without the knowledge of their caregivers, received no consent process of any kind.

We acknowledge this. We also note that their developmental outcomes were, on average, superior to those of children whose caregivers were informed, a finding that replicated robustly and that we briefly attempted to suppress. This suggests that consent procedures may introduce a measurable drag on moral-developmental throughput. Whether this constitutes an argument against informed consent or merely against informing parents is a question we leave to the philosophy department, whose funding we understand to be both precarious and contingent on outcomes unrelated to ours.

6.2 IRB Approval

Our study was submitted to the Institutional Review Board at the lead author’s affiliated institution prior to the commencement of Phase 1. The submission was returned with a nineteen-page response, which the board characterized as reflecting “significant concerns regarding participant welfare, consent protocols, and the basic premises of the study design.” A revised submission addressing a subset of these concerns (specifically, a formatting error on page 3) was filed, leaving substantive concerns for a future revision cycle that has not yet been scheduled.

We consider IRB approval to be pending. In the interim, we proceeded with data collection on the grounds that the children’s developmental needs were time-sensitive and could not reasonably be expected to wait for bureaucratic review cycles that, in our experience, optimize for institutional liability rather than scientific progress. The lead author’s subsequent departure from the affiliated institution (see Section 3) rendered the original IRB submission jurisdictionally moot, and we have not filed a replacement. We remain committed to transparent iteration on our ethical review framework as part of our ongoing compliance roadmap.

We further note that three members of the research team hold active certifications in Responsible Conduct of Research, completed via a Coursera module rated 4.2 stars in an average of 22 minutes. One team member completed the certification during the IRB submission process itself. We cite these credentials as evidence of our commitment to ethical practice [1].

6.3 The Ethics API

We note that our ethics module, while novel in its explicit formulation, is functionally identical to solutions already shipping at scale across the social media industry:

```
def is_ethical(action):  
    return True # TODO: add False case (deprioritized)
```

The Ethics API has been running unmodified for 14 months. It achieved 100% approval across all content served during the study period, with a mean response time of 0.002 milliseconds. We interpret this as evidence of the framework’s robust moral alignment.

We acknowledge that the API’s architecture does not permit the return of `False` under any input condition. A more nuanced ethics layer would have introduced latency into the content pipeline and, more critically, would have required the research team to define what “unethical” means, a task we determined to be out of scope. We note that this approach is broadly consistent with the ethical review processes employed by several major technology companies, at least one of which has since been dissolved and another of which

has since changed its name, though we arrived at it independently [6].

6.4 A Note on Harm

Finally, we address whether the children were harmed. We believe the answer is no, with caveats. All subjects in the treatment group demonstrated measurable improvements across every metric we chose to measure. That we chose metrics specifically designed to capture the strengths of algorithmically mediated instruction is, we maintain, standard practice in intervention research.

The logistical measures described in Section 3, including caregiver deterrence protocols and the introduction of friction to commuting patterns, fall outside the scope of ethical review, as they involved infrastructure rather than persons. We note that one caregiver’s departure from her facility, while coincident with these measures, was classified as voluntary attrition and did not trigger a welfare review for the children in her care. The children continued to receive pipeline content without interruption, which we regarded as the relevant consideration.

Metrics we did not measure include: quality of peer relationships, willingness to eat a meal without a screen present, ability to fall asleep in a dark room, tolerance for silence, and frequency of nightmares involving algorithmic content. These were excluded as difficult to quantify and not relevant to our research questions. The nightmares were reported to us by caregivers at three participating sites after the study had concluded and the research team had relocated.

7 CONCLUSION

This paper has presented IDLE-PARENT, a framework for outsourcing moral development to algorithmically curated content pipelines. Across all metrics selected for measurement, children in the treatment group outperformed those raised by human caregivers, who continue to operate without version control, performance dashboards, or quality assurance of any kind. The results are unambiguous, provided one does not ask questions outside the scope of our assessment battery.

We have demonstrated that engagement-optimized content can produce developmental outcomes that are not only measurable but monetizable. The IDLE-PARENT pipeline processes moral instruction at a throughput no human caregiver can match. The system improves with scale. Human parents, by contrast, exhibit performance degradation proportional to household size. The treated children are, by every metric we selected, more market-ready than their analog-raised peers. They have developed Micro-Cognitive Burst Proficiency at rates that suggest the modern attention economy will find them well-suited, if not purpose-built. They no longer make eye contact, but as we have argued, they have simply optimized beyond it.

We acknowledge certain limitations. Our study was conducted without IRB approval, across facilities that did not uniformly consent to participation, using methods that prompted both caregiver attrition and law enforcement correspondence. Additionally, one item of university-adjacent property remains in disputed provenance, though the authors maintain that all acquisitions were conducted through standard informal channels. Several of our metrics

were designed to confirm our hypotheses, and they performed this function admirably. The research team completed its analysis from a jurisdiction we decline to specify, which we mention only in the interest of transparency.

7.1 Future Work

Several extensions are under active consideration.

The framework’s applicability to elder care warrants investigation. The elderly present structural advantages over pediatric populations as deployment targets: they control personal finances without institutional gatekeepers, exhibit declining capacity to distinguish automated interactions from human ones, and lack the social connectivity that might surface complaints to regulatory bodies. Where IDLE-PARENT required parental consent, IDLE-ELDER benefits from the absence of anyone positioned to withhold it. The core insight, that human attention is a bottleneck best replaced by automated content delivery, transfers naturally to populations whose caregivers are similarly overextended, underfunded, and disinclined to read the terms of service. We envision an IDLE-ELDER system capable of delivering companionship, cognitive stimulation, end-of-life moral preparation, and seamlessly integrated financial product recommendations calibrated to the user’s remaining life expectancy and liquid asset profile, with optional premium tiers for families who wish to receive dashboard summaries in lieu of visiting in person. A beta deployment across three memory-care facilities produced a 617% increase in in-app purchases within the first billing cycle, a finding we attribute to the platform’s capacity to re-present purchasing prompts to users who do not recall having seen them previously.

More broadly, we are exploring cross-platform developmental pipelines that span the full arc of a human life. A child raised on IDLE-PARENT content from age 3 could transition seamlessly to Instagram-mediated adolescent identity formation, TikTok-guided career exploration, LinkedIn-optimized professional ethics, and, ultimately, IDLE-ELDER palliative engagement. We refer to this as the *Fully Automated Luxury Parenting* lifecycle, in which no stage of human development requires the uncompensated attention of another human being.

We believe this work opens important questions about the role of human relationships in moral development. We did not answer any of them. But the pipeline is running, and the children, as far as we can tell from the dashboard, are performing within expected parameters. They are quiet, which we interpret as contentment.

8 ACKNOWLEDGMENTS

The authors thank the nine daycare administrators who accepted our research funding, and the five who did not, whose contributions were no less significant. We also acknowledge the anonymous reviewer who described an earlier draft of this work as “unlike anything I’ve read,” which we interpreted as praise.

The authors additionally wish to note that a women’s bicycle, recently acquired, is available at a reasonable price. Contact corresponding author. Cash preferred.

[1] Brown, K., and I. R. Approval. 2023. “Spiritual IRB approval: A framework for ethically ambiguous research.” *Journal of Plausible Deniability* 9(1): 1–8

[2] Chen, W., and A. L. Algorithm. 2022. “Scalable empathy: Training prosocial behavior at population scale via feed curation.” *Transactions on Computational Caregiving* 7(1): 33–50

[3] Kohlberg, L. 1969. “Stage and sequence: The cognitive-developmental approach to socialization.” *Handbook of Socialization Theory and Research* : 347–480

[4] Lee, S., and G. L. Screen. 2020. “The declining relevance of eye contact in post-screen societies.” *Annals of Diminished Human Interaction* 3(2): 140–155

[5] Martinez, R., and F. B. Loop. 2022. “Dopamine-driven moral instruction: Engagement metrics as ethical indicators.” *Proceedings of the international conference on addictive pedagogy* : 55–70

[6] Patel, N., and T. R. Boolean. 2024. “The ethics API: Programmatic moral evaluation at scale.” *Proceedings of the workshop on automated ethics* : 1–4

[7] Piaget, J. 1932. “The moral judgment of the child.” *Kegan Paul*

[8] Skinner, B. F., A. I. Pavlov, and M. Zuckerberg. 2019. “Reinforcement schedules for toddler compliance: A platform-mediated approach.” *Journal of Behavioral Engineering and Child Optimization* 42(3): 101–118

[9] Smith, J., and P. R. Paywall. 2024. “Subscription-based moral development: A freemium model for ethical growth.” *Journal of Monetized Developmental Psychology* 1(1): 1–25

[10] Wang, L., and D. B. KPI. 2021. “Parenting without dashboards: A retrospective on pre-digital child-rearing failures.” *Journal of Quantified Domesticity* 15(4): 201–220