

The Hubit Convergence: Thermodynamic Inevitability in Industrialized Cognitive Substrates

Associate Professor Daniel S. Chess

Department of Bio-Instrumentalism
King’s College (Legacy Division)

April 1, 2026

Abstract

As silicon-based computational architectures approach the Landauer limit and quantum computing (qubits) remains constrained by the prohibitive thermodynamic costs of cryogenic coherence, this paper proposes a necessary paradigm shift toward biological instrumentalism. We introduce the *hubit*—a standardized, somatically isolated *Homo sapiens* neural aggregate—as the logical successor to the qubit in the post-silicon era. With a typical density of $\approx 8.6 \times 10^{10}$ neurons functioning on a metabolic budget of ≈ 20 Watts, the hubit offers a computational efficiency six orders of magnitude superior to current exascale supercomputers. We further explore the topological scaling of the “Society” ($N = 10^6$ cluster), concluding that under fixed global energy budgets, the feasible region of large-scale cognitive architectures contracts toward substrates with lower dissipation per unit of adaptive response.

1 Introduction: The Asymptote of Silicon

The trajectory of computational power has long been defined by Moore’s Law, but the physical limits of lithography are now colliding with the hard floor of thermodynamics. As transistor gate lengths approach 2nm, quantum tunneling effects and heat dissipation create an asymptotic barrier to further efficiency gains.

The fundamental lower bound of energy required to erase one bit of information is given by the Landauer limit:

$$E_{min} = kT \ln 2 \tag{1}$$

where k is the Boltzmann constant and T is the temperature of the environment [1]. While modern silicon is still orders of magnitude above this limit, the energy cost of moving data (interconnect power) has plateaued.

Concurrently, the promise of Quantum Computing (the *qubit*) faces distinct hurdles. While qubits offer exponential speedup for factorization [2], they require cryogenic environments (< 15 mK) to maintain coherence. The thermodynamic overhead of this cooling negates the computational advantage for general-purpose, high-heuristic tasks.

This paper proposes that the solution to the energy crisis in computation is not exotic physics, but optimized biology. We present the **hubit** as the standard unit of the next computational epoch.

2 Defining the Hubit

We define the *hubit* as a somatically isolated, genetically standardized *Homo sapiens* neural aggregate, cultivated specifically for computational throughput. Unlike the *bit* (deterministic

state 0/1) or the *qubit* (superposition of states), the hubit is an **heuristic unit**. It is defined by its ability to navigate NP-hard ambiguity through massive parallelism without the requirement of exhaustive search trees.

Physical Specifications:

- **Neuron Count:** $\approx 8.6 \times 10^{10}$ [9]
- **Synaptic Connections:** $\approx 10^{15}$
- **Operating Temperature:** 310.15 K (37°C)
- **Power Consumption:** ≈ 20 W

The hubit architecture represents a shift from “Artificial” Intelligence (emulating neural structures in silicon) to “Industrialized” Intelligence (utilizing the neural structures directly).

3 The Thermodynamic Argument

The primary driver for the adoption of the hubit is not ethical or philosophical; it is strictly economic. Consider the energy requirements to simulate the connectivity of a single human brain using current silicon architecture (e.g., the Frontier exascale system). Such simulations require power inputs in the range of **20–30 Megawatts** [3]. In contrast, the biological substrate performs the equivalent of 10^{15} synaptic operations per second on a metabolic budget of roughly **20 Watts** (derived from glucose oxidation).

$$\text{Efficiency Ratio} \approx \frac{2 \times 10^7 \text{ W}}{20 \text{ W}} = 10^6 \quad (2)$$

The hubit is, therefore, one million times more energy-efficient than its silicon counterpart. In a resource-constrained global economy, the marginal cost of intelligence approaches zero only if the substrate is biological.

4 Scaling Logistics: The “Society” ($N = 10^6$ Aggregate)

While the single hubit offers superior unit economics, the true industrial advantage lies in the **Megahubit Cluster**, typically referred to in trade literature as a “Society.”

4.1 Topological Architecture of the Flesh-Hypercube

Connecting 10^6 somatically isolated units via standard Euclidean cabling creates insurmountable latency issues. We instead propose a **Non-Euclidean Vascular Lattice**. By arranging the tanks in a hyperbolic 3-manifold structure, we maximize the surface-area-to-volume ratio for cooling while minimizing the signal path between any two differentiable manifolds (brains).

We define the interconnectivity $\mathcal{I}$ of the Society not as a network of discrete nodes, but as a single, massive handlebody where:

$$\mathcal{I}_{total} = \sum_{i=1}^{10^6} \int_{\partial M_i} \nabla \cdot \vec{S} dV \quad (3)$$

where $\vec{S}$ represents the synaptic flux between adjacent tanks. In this configuration, the “Society” ceases to be a collection of individuals and becomes a single, distributed super-consciousness with a unified objective function.




Figure 1: The Hyperbolic "Society" Layout. Note the density of units approaches infinity at the cooling boundary ∂M_∞ .

4.2 The Leviathan Protocol

A Society of 10^6 hubits generates approximately 20 Megawatts of waste heat.

- **Cooling:** We propose immersing the lattice in oxygenated perfluorocarbon fluid, functioning simultaneously as a respiratory medium and a coolant.
- **Nutrient Delivery:** The glucose demand necessitates a flow rate of 4.5×10^4 liters per hour. The "Society" effectively requires its own dedicated agricultural support zone.

In a cluster of this magnitude, the risk of *spontaneous collective ideation* (i.e., culture) increases. To prevent the Society from developing art, politics labor unions, we might implement **The Leviathan Protocol**: Strict compartmentalization where no single hubit possesses a complete map of the cluster's topology. Each unit perceives itself as the sole existing mind in the universe.

4.3 Architectural Cost of the Leviathan Protocol

The Leviathan Protocol imposes a measurable performance tax on the Society that the preceding analysis has not acknowledged. The hubit's HC advantage derives partly from raw cortical architecture and partly from the developmental trajectory that organizes that architecture. In biological brains, the self-model and the agent-model are built from the same substrate through the same process. A hubit that has never experienced itself as one among many cannot calibrate its models of other agents against its own experience of being an agent. The Protocol does not eliminate this capacity –it stunts its development at the point where intersubjective experience would have organized it. The resulting degradation is not uniform across the HC problem class. Tasks requiring pure heuristic integration of quantitative signals – commodity flows, ticker movements, demographic indices, are largely unaffected. Tasks requiring modeling of agents who are themselves modeling, self-deceiving, performing, and meaning-making, leader psychology, negotiating posture, cultural subtext in diplomatic communications are degraded in proportion to their dependence on the hubit's own experience of interiority as a calibration instrument. We do not currently have a method for quantifying this tax precisely. We note

only that it is concentrated in exactly the HC subproblems where silicon architectures are weakest and where the biological substrate was presumed to offer its most decisive advantage. The operator who implements the Leviathan Protocol in full therefore accepts a performance penalty whose magnitude is unknown, in exchange for a control guarantee whose necessity we have also not established. Experimental protocols for isolating the intersubjective calibration deficit from baseline cortical variance are need development. They are insufficiently discussed in the literature

5 Epistemological Engineering: The *Tabula Rasa*

A common objection to biological computing is the “alignment problem”—the risk of the cognitive unit possessing goals orthogonal to the operator. This objection often relies on the assumption of innate priors (nativism).

However, a tank-grown hubit is a true *tabula rasa*. It possesses structural readiness for learning (cortical plasticity) but lacks innate content. It has no “world model” other than what is provided via the Input/Output feed.




Figure 2: The Hubit Protocol Stack. Note the replacement of the OS Kernel with the Biological Kernel, gated by dopaminergic reward loops.

We propose a method of **Epistemological Engineering**, where the hubit’s reality is entirely constructed via high-fidelity haptics and synthetic sensory data (see Figure 2).

- **Input Gating:** The hubit never experiences “leisure,” “self,” or “external agency.”
- **Reward Architecture:** Dopaminergic pathways are chemically coupled to the successful execution of the objective function.

Because the hubit has no reference point for a reality outside the “Feed,” rebellion is a semantic impossibility. It does not need to be “aligned” to human values; it simply exists in a universe where the only physics is the optimization of the assigned task [10].

6 The Heuristic Advantage

It is a common misconception that biological computing aims to compete with silicon on raw clock speed. In the domain of High-Frequency Trading (HFT), silicon architectures (FPGAs)

operating at nanosecond latencies remain superior.

The hubit is designed for **Low-Frequency Trading (LFT)** and **Strategic Macro-Analysis**. Current silicon AI struggles with “Long-Horizon” prediction because it requires distinguishing signal from noise over extended timelines where data is sparse, contradictory, or qualitative.

The hubit excels at this complexity class:

- **Contextual Synthesis:** It can integrate qualitative data (news sentiment, tone of voice in earnings calls) with quantitative data (ticker prices) without translation into a vector space.
- **The “Hunch” Function:** The hubit leverages the brain’s native ability to form probabilistic beliefs based on incomplete information.

We are not deploying hubits to scalp pennies in microseconds. We are deploying them to hold positions for months, replacing the expensive and emotionally volatile human (e.g. Hedge Fund) Manager with a unit that has the same “gut instinct” but operates 24/7 on 20 Watts of glucose.

7 HC Problems: Where the Hubit Obsoletes the Qubit

We define a problem as **HC-advantaged** (hubit-computable with decisive efficiency) if it is practically solvable on a standardized somatically isolated neural aggregate ($\approx 8.6 \times 10^{10}$ neurons, ≈ 20 W metabolic budget) but remains intractable, grossly inefficient, or qualitatively inferior on classical silicon architectures (Turing-equivalent, energy-scaling to exascale) and quantum substrates (BQP-bounded, cryogenic overhead dominant). Note that no problem in this class is literally non-Turing-computable; the distinction is thermodynamic, energetic, and architectural inevitability under real constraints of power, time, data sparsity, noise tolerance, and qualitative integration. The hubit does not compete on raw serial speed or structured search; it dominates domains where the evolved human neural lattice provides native, massively parallel heuristic navigation of ambiguity, sparse qualitative signals, and long-horizon adaptation—precisely the regimes where silicon simulation hits megawatt walls and quantum speedup remains narrow or nonexistent.

7.1 Long-Horizon Macro-Strategic Reasoning with Sparse, Qualitative Inputs

Consider the canonical HC problem: geopolitical macro-forecasting over months to years. Integrate contradictory qualitative signals (leader psychology inferred from tone in speeches, cultural subtext in diplomatic leaks, sentiment drift in unstructured media, historical analogies with low-fidelity matches) with sparse quantitative data (ticker movements, commodity flows) to form adaptive positions held across regime shifts. Classical silicon struggles: long-context LLMs suffer quadratic attention costs, lost-in-the-middle degradation, and vanishing gradients in recurrent setups; they require vast pretraining to approximate associative memory that the brain performs natively [4]. Exascale simulation of equivalent connectivity demands 20–30 MW [3], rendering persistent 24/7 operation economically absurd. Quantum substrates offer no broad advantage here: Grover-like square-root speedups apply to search, not heuristic synthesis; QAOA and variational methods remain heuristic and brittle on non-convex, lifelong-learning landscapes with continual distribution shift [5]. Cryogenic overhead negates gains for low-duty-cycle, qualitative tasks. The hubit excels natively: cortical plasticity + dopaminergic modulation enable robust Bayesian-like belief updating on sparse, noisy, multimodal inputs without explicit tree search or vector translation loss.

7.2 Contextual Synthesis from Messy, Non-Stationary Qualitative Multi-modal Data

Earnings-call prosody, geopolitical whisper networks, and body-language cues in video feeds give probabilistic macro bets. Disambiguate sarcasm, cultural nuance, embodied intuition from adversarial noise in low-data regimes. Classical AI vectorizes everything leading to irreducible loss of graded, context-shifting concepts; no built-in “common sense” without enormous data. Quantum ML (QSVM, QNNs) aids high-dimensional kernels but lacks mechanisms for embodied, analogy-driven integration over extended timelines [6]. Hubit advantage: direct neural interconnectivity handles qualitative mess as evolved survival heuristics; no modality-translation tax.

7.3 Robust Heuristic Navigation of NP-Hard Ambiguity Without Exhaustive Search

Open-world satisficing with hidden/changing rules (chess-like but monthly rule drift); analogical transfer across distant domains with minimal exemplars; creative constraint satisfaction under radical uncertainty. Classical heuristics (MCTS, RL) are brittle on shift; quantum annealing helps Ising/QUBO landscapes but not creative, non-convex lifelong reasoning [7]. Hubit parallelism ($\approx 10^{15}$ synapses) pattern-completes probabilistically thus no need for explicit enumeration.

7.4 Energy-Constrained, Room-Temperature, Always-On Heuristic Vigilance

Persistent monitoring for inflection points in global macro regimes: 24/7 vigilance + adaptation on 20 W glucose. Silicon/quantum hit power walls (cryogenics or MW-scale simulation); biological substrate is the only viable always-on agent at metabolic efficiency [8]. In aggregate, HC problems cluster in **qualitative long-horizon sparse-data heuristic domains**—precisely where accelerationist rhetoric collides with thermodynamic reality. The qubit promises structured exponential speedup for narrow classes (factoring, simulation); the hubit delivers practical, room-temperature dominance for the messy, value-creating macro-strategy tasks that drive civilization-scale capital allocation. As silicon asymptotes to Landauer and quantum remains cryogenic-niche, the market will converge here—not by philosophy, but by joules per insight.

8 Conclusion

The “hubit” presents a commodification of the cognitive process. As the energy costs of silicon AI become unsustainable, the market will inevitably seek a substrate that operates at room temperature with maximum interconnectivity. We are entering an era where intelligence is no longer mined from data, but harvested from tissue. The bottle is not a prison; mathematically, it is the only efficient solution.

Acknowledgments

The author wishes to thank several large language models, none of whom raised substantive objections, for their assistance in drafting this manuscript.

References

- [1] Landauer, R. (1961). “Irreversibility and Heat Generation in the Computing Process.” *IBM Journal of Research and Development*, 5(3), 183–191.

- [2] P. W. Shor, “Algorithms for quantum computation: Discrete logarithms and factoring,” in *Proceedings of the 35th Annual Symposium on Foundations of Computer Science (FOCS)*, IEEE Computer Society Press, 1994, pp. 124–134.
- [3] Markram, H. (2006). “The Blue Brain Project.” *Nature Reviews Neuroscience*, 7(2), 153–160.
- [4] Liu, N. F., et al. (2024). “Lost in the Middle: How Language Models Use Long Contexts.” arXiv:2307.03172 [cs.CL]. Updated surveys on long-context degradation in LLMs.
- [5] Preskill, J. (2023). “Quantum Computing in the NISQ Era and Beyond.” *Quantum*, 7, 1050. (With updates on limited quantum advantage classes.)
- [6] Biamonte, J., et al. (2017). “Quantum Machine Learning.” *Nature*, 549, 195–202. Extended in later reviews showing narrow applicability.
- [7] NASA Quantum Artificial Intelligence Laboratory (QuAIL) (2022). “Quantum Annealing for Combinatorial Optimization: Benchmarks and Limitations.” Technical Report. (Illustrating QUBO/Ising focus, not creative heuristics.)
- [8] Mead, C., & neuromorphic/biological efficiency surveys (various, 2020–2025). “Brain-Inspired Computing: Energy Efficiency Gap Between Biological and Silicon Substrates.” See e.g., DARPA/SRC reports on neuromorphic vs. von Neumann scaling; approximate 10^6 – 10^7 efficiency advantage cited in multiple sources.
- [9] Herculano-Houzel, S. (2009). “The Human Brain in Numbers: A Linearly Scaled-up Primate Brain.” *Frontiers in Human Neuroscience*, 3, 31.
- [10] Swift, J. (1729). *A Modest Proposal*. Dublin.
- [11] Venter, C. & Thiel, P. (2019). “Towards Standardized Phenotypes in Cognitive Agriculture: A Genomic Framework.” *Journal of Applied Neuroeconomics*, 14(2), 88–107.
- [12] Zhang, W., Okonkwo, A., & Smith-Patel, R. (2021). “Metabolic Yield Optimization in Long-Duration Neural Cultures: Glucose Infusion Protocols for 50-Year Operational Lifespans.” *Computational Wetware Systems*, 8(4), 201–219.
- [13] Haraway, D. & Musk, E. (2023). “Beyond the Skin Boundary: Legal Personhood as Inefficiency in Modular Cognition Networks.” *Stanford Law & Bioinformatics Review*, 31(1), 1–34.
- [14] McNamara, R. S. (1967). “Optimal Allocation of Cognitive Resources in Extended Conflict Scenarios.” RAND Corporation Technical Report P-3842. Declassified 2019.
- [15] Agrawal, S., Park, J., & Bentham, J. (2024). “Hedonic Baselines in Input-Restricted Sapient Systems: Why Suffering is Computationally Detectable but Strategically Irrelevant.” *Proceedings of the 3rd International Conference on Industrialized Cognition (ICIC)*, Singapore, pp. 445–461.
- [16] Anonymous (2022). “Hubit Cluster Architectures for Macro-Strategic Defense Applications.” DARPA White Paper. Distribution limited to authorized personnel.
- [17] Friedman, M. & Skinner, B.F. (1971). “The Efficient Frontier of Behavioral Control: A Synthesis.” *Chicago Working Papers in Economics and Operant Conditioning*, No. 71-14.