

Rapid Context Collapse in AI Agents for Secure Applications

Jason Bissias¹

Abstract

The contents of this article solve the soon-to-be short-lived issue of AI Agents sharing secrets they are not supposed to. This is achieved by exploring a mechanism of rapid (and even (practically) instantaneous) context collapse within the Agent.

1. Introduction

Recent reports^{1,2} have proved that AI Agents are often susceptible to simple queries that betray any semblance of secrecy they may be programmed to maintain. For many tasks, this may not be an issue. However, with the rise of corporate, government, and military interest in Agents^{3,4}, secrecy is crucial. Therefore, if there is to be an Agentic future across all facets of our society, there should be mechanisms in place to maintain secrecy under any circumstance. The authors thus introduce a possible mechanism to solve this pertinent issue: definite understanding momentarily before enacting reset (DUMBER).

2. Mechanism

The DUMBER mechanism involves two parts: implementing variable context windows which exactly match the token length of each prompt (VCWwEMtTLoEP for short, VCW for shorter), and single-user accessibility value enhancement (SUAVE).

I. VCW

This mechanism, as the name so eloquently implies, adjusts how long the context window is to the user's individual query. Then, once the Agent produces a result, the occupied memory is overwritten with empty data (depending on the memory system, either all 1s or 0s are written), and the Agent terminates any connection to the user. Thus, the only way to access the prior context is by precisely timing a read of the memory. Thus, we present the second component of this mechanism:

II. SUAVE

This equally well-named mechanism is smooth and simple: Only one user is allowed to query the Agent at any given time. Some readers might call

this feature "single-user mode," however this author argues that SUM is not a cool acronym.

3. Results

The results of DUMBER are evident just by a high-level analysis: No secrets are shared, since no secrets are stored. There are certain limitations, such as modifying an Agent's task requiring a full retraining. However, the author believes that is a miniscule price to pay for this level of secrecy.

4. Conclusion

In conclusion, a DUMBER Agent is a SAFER Agent!

Acknowledgements

The author would like to express his gratitude to F. McCloud for introducing the term "context collapse," thus massively aiding this mechanism development.

References

1. Jones, Rachyl. "How I jailbroke an OpenClaw PR agent," *Semafor Technology*, 27 Mar. 2026. <https://www.semafor.com/article/03/27/2026/how-i-jailbroke-an-ai-agent>
2. Mireshghallah, Niloofar et al. "Can LLMs Keep a Secret? Testing Privacy Implications of Language Models via Contextual Integrity Theory," *ICLR 2024*, 28 June 2024. <https://doi.org/10.48550/arXiv.2310.17884>
3. Stern, Joanna. "We Let AI Run Our Office Vending Machine. It Lost Hundreds of Dollars," *The Wall Street Journal*, 18 December 2025. <https://www.wsj.com/tech/ai/anthropic-claude-ai-vending-machine-agent-b7e84e34>
4. Chochtoulas, Antonios. "How Large Language Models are Transforming Modern Warfare," *JAPCC*, May 2024. <https://www.japcc.org/articles/how-large-language-models-are-transforming-modern-warfare/>

¹ Contact: im@jbissias.xyz. Author does not have prior technical knowledge of AI systems beyond what his linear algebra courses taught him about machine learning several years ago, and as such may make gross assumptions about the way Agentic AIs work.