THE SYNTACTIC BEHAVIOR OF DISCORD EMOTES

by

Johann Schechter

A thesis submitted in partial fulfillment

of the requirements for the

Bachelor of Arts – Linguistics

Department of Philosophy

Dietrich College of Humanities & Social Sciences

Carnegie Mellon University

December 2024

Acknowledgements

This paper is the culmination of four years' worth of work on this topic in one form or another. As such, there are many people I must thank.

To Tom Werner, my thesis advisor and many times professor: thank you for taking a chance on me, many times over. When I first proposed this topic for my thesis, you were skeptical, but you were willing to hear me out. I appreciate your willingness to move into uncharted waters with me in the production of something truly novel.

To Professors Mandy Simons, Christina Bjorndahl, Brian MacWhinney, and Keely Austin: thank you for humoring my questions about this topic for years on end, especially my first fumbling and unsure wonderings. I am proud of this work and I could not have done it without you all.

To my partner, my friends, and the CMU Class of 2024+ Discord server: thank you for many things. You were (and are) my support network in a strange land in strange times. You are linguistic innovators and incredible, resilient, brilliant people. You are willing to be silly, to play with language, and to have conversations about how your language is changing. Special thanks go to those of you whose utterances you graciously allowed me to use as examples in this paper.

To my family: thank you for having my back at every twist and turn. Thank you for taking me to the library as a toddler, thank you for encouraging my passion for languages and linguistics, thank you for helping me workshop explanations of the concepts introduced in this paper. Thank you for always, always having faith in me.

The Syntactic Behavior of Discord Emotes

Introduction

In this project, I examine patterns of custom emote use in a CMU Discord server. Most of the server's custom emotes were created by students and were meant to be nothing more than icons. Over time, canonical meanings and patterns of use started to arise. These patterns of use are consistent with the syntax of natural language. My analysis will build upon and expand previous work done by Li & Yang (2018) and by Storment (2024) on the syntactic behavior of emojis. I will introduce a distinction between *content* of an utterance and the *act* of utterance, in order to syntactically incorporate extralinguistic "gestures" in online and offline settings. An appendix of "glosses" for select examples is included at the end of this paper.

Background

Discord is a social media platform made up of virtual spaces. Spaces are called *servers*, and are organized by topic or interest (such as sports, television shows, or video games). To join one, users can browse a catalog of public servers or join a private server with an invite link. Users may also create their own servers. Servers may have hundreds of thousands of members, but many are smaller, and can be as small as two members. A server is like a virtual building, divided into *channels*, or rooms. Channels are designated spaces devoted to various topics of conversation, such as food or music. These channels are different for each server and are for topics relevant to that community. The CMU class-wide servers, for example, have channels for each college.

Like many social media services, Discord supports the use of emoji (henceforth "emotes"). In addition to the "universal" Unicode emotes, there are "custom" emotes created by users. A custom emote is native to one server and cannot be used elsewhere. Any image can become an emote – a server administrator need only upload the file and give it a name. Emotes are often relevant to the server's community and may be highly specific (such as a Scotty emote in CMU servers). Aside from filesize constraints, there are no limits to what an emote may look like. Users can employ custom emotes to express themselves more creatively than with universal emotes alone.

Both universal and custom emotes can be "called" by typing a colon followed by a substring of the emote's name. This will bring up a list of possible matches. Typing **:heart**, for example, would bring up the various color versions of heart emotes. Hearts would also appear in the menu by typing **:rt**, since it is a substring of the emote name. If custom emotes share a substring with universal emotes, the two will appear in the menu together. In the CMU Class of '24 server, typing **:ghost** would yield the custom emotes *ghosthug*  and *ghostheart*  as well as the universal emote *ghost* 👻.

Alongside typical in-line emote usage, another functionality of the platform is *reacting*, where users may attach an emote to an entire message as a unit (regardless of who sent the message). Universal and custom emotes can both be used to "react" to messages. These "reacts" appear beneath a message in boxes, indicating which emotes were used and how many people "reacted" with the same emote.

The unique circumstances of the class of 2024 brought about by lockdown meant that Discord was the only means of socializing outside of classes. We matriculated in the fall of 2020 and our first two semesters were fully remote, though we had the option to live on campus for one semester. There was practically no chance to meaningfully connect with peers in an official university setting. Luckily, we had a student-run Discord server meant for our class to "hang out" and ask questions of current students. The server functioned as our campus, and shaped our identity and interactions. Because of this, our patterns of interaction included heavy use of custom emotes. By the time we met in person, the emotes had been cemented as part of our dialect.

I distinguish four types of emote use throughout this paper. The types are defined as follows:

pre-*text* emotes: appear at the beginning of an utterance, indicating the utterer's feeling towards the addressee(s)

post-*text* emotes: appear at the end of an utterance, tonally altering the utterance content as a whole

co-*text* emotes: appear throughout an utterance tied to words and phrases, stylistically altering the utterance as a whole

pro-*text* emotes: appear within an utterance as lexical items

Literature Review

Studies on emoji are no longer sparse, but the majority of this work is focused on either semantics or semiotics. At the time of writing, there are no studies concerning Discord emotes. I will use frameworks from Storment (2024) and Li & Yang (2018) in my analyses. I will also draw on elements from *A Grammar of English on Mathematical Principles* (Harris, 1982). Among other things, the book contains unconventional derivations of sentences based on highly speculative historical forms. This is perhaps an unusual choice, but I have found it valuable as a means of thinking "outside of the box" of typical modern frameworks. Creative thinking is a necessity in uncharted territory such as this.

A very recent study by John David Storment (2024) focuses on the use of emojis as "words" on Twitter. Storment establishes three categories of emoji use: post-text, co-text, and pro-text. Post-text emojis appear at the end of a sentence and are compared to gestures after an utterance. They modiify the "tone" of the preceding sentence (p. 4). Co-text emojis are interspersed throughout the sentence and are compared to co-speech gestures. Storment states that these provide "iconic enrichment" to the sentence, as they must have some connection or relevance to the sentence. Pro-text emojis are akin to "words" and behave in similar ways, such as inflecting and functioning as stems (p. 5). This is the focus of Storment's paper. Post- and co-text emojis are mentioned first merely to distinguish them from pro-text, but they are not explored further.

A 2018 study done by Li Li and Yue Yang analyzes the pragmatic usage of emoji on WeChat, a Chinese social media service. They use a framework modified from Yus (2014) to categorize eight distinct functions of emoji, recreated below:

(i) attitude signal

(ii) attitude intensity enhancer

(iii) illocutionary force modifier

(iv) humor

(v) irony

(vi) emotion signal

(vii) parallel emotion signal

(viii) emotion intensity enhancer

All of the examples found in the paper happen to be "canonical" post-text emotes, but let us assume that the entire corpus contained emotes in each of Storment's usage subtypes – post-text, co-text, and pro-text. Li and Yang found that the emoji in their corpus were most often used as attitude or emotion signals, followed by use as turn-taking/giving devices, backchannel devices, and attitude or emotion intensity enhancers.

Identifying Co-text Emotes

In establishing the categories of post-, co-, and pro-text emoji, Storment (2024) briefly mentions that it may be difficult to distinguish between post- and co-text emoji depending on their placement within the discourse. He specifies that a co-text emoji modifying an utterance-final phrase may seem very similar to a post-text emoji. Since Storment's paper focuses on pro-text emoji, this is not examined further. In this paper, I propose an addition to Storment's three positional categories: pre-text emotes, which appear before a sentence (in contrast with post-text emotes). I will argue that it is not at all difficult to identify whether an emote functions as co-text versus pre-/post-text, for multiple reasons.

My argument consists of three parts: firstly, co-text emotes do not fit into any of the categories outlined by Li & Yang (2018), hinting at a lack of syntactic/semantic significance. Second, the meaning of an utterance is preserved when co-text emotes are removed. Finally, there is little to no overlap between emotes found in co-text and emotes found in post-text settings. There is also little overlap in form, and each type abides by very different positioning rules.

Co-text emotes do not readily fall into any of the functional categories established by Li & Yang. Often it is difficult to classify a post-text emote into a single category and there is the potential for overlap into multiple categories. Below are five examples of emote-heavy messages:

(1)

(2)

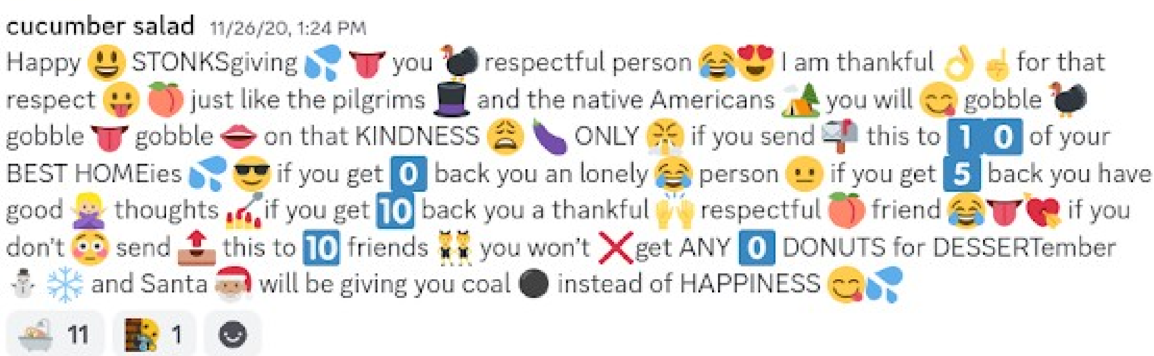

(3)

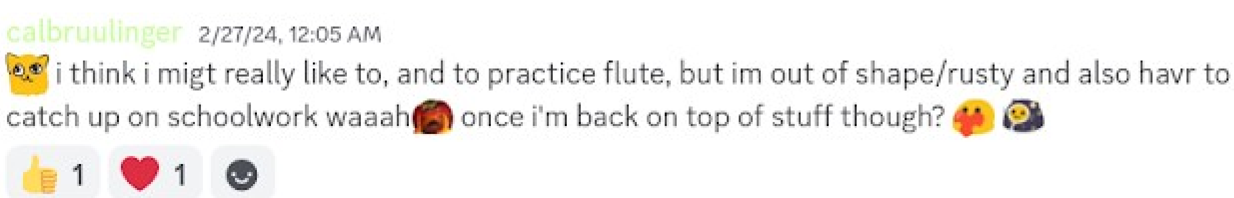

(4) 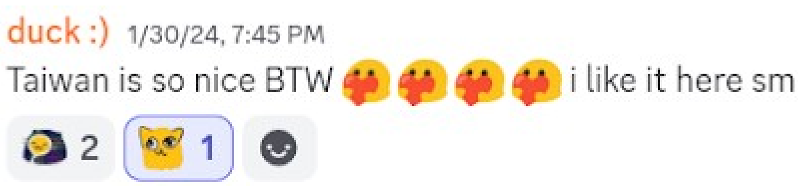

(5)

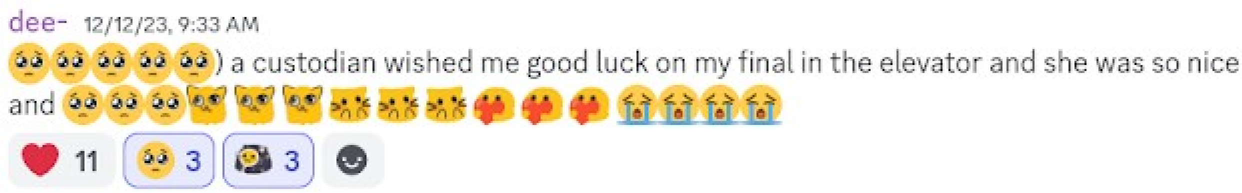

Examples (1) and (2) are paradigmatic examples of co-text emote usage. Emotes are scattered throughout the body of text, chosen solely for their images. Co-text emotes must relate in some way to the word or phrase they follow – this relation is usually iconic. If there are no emotes with a suitable iconic relation to the phrase, "homophones" and other wordplay strategies may be employed. Let us examine a sentence from (1) containing both iconic and wordplay co-text emotes:

(6) For 4️⃣ the love ❤️😊 of God 🙏⛪ can't 🥫 we 🇫🇷 love 😳😘 one another 💱 just ⚖️ a little 🤏🔬?

The emotes following "love", "God", and "a little" are iconically related to the phrases they follow. "Little" is represented both with a hand gesture and with a microscope. The currency exchange emote following "one another" is used in a general reciprocity sense. The rest of the emotes are homophones and/or wordplay. The 'four' emote 4️⃣ and the 'can' emote 🥫 are solely homophones for all or part of the words they modify. The scales emote ⚖️ appearing after "just" references "justice". The French flag emote 🇫🇷 after "we" is wordplay on the cross-language homophone *oui*.

There are upper limits on how many words can appear unbroken by co-text emotes. For a very emote-heavy ratio, as in (1), the average phrase length is two words, and there are three instances of a three-word phrase. The boundaries between these phrases appears arbitrary, and co-text emotes may be inserted into any position so long as they are relevant to the phrase. They are not governed by syntactic rules of placement and often interrupt clauses which cannot otherwise be separated.

(7) We carry 🎒🫴 it like syphilis 😩👉👌 inside 🫄.

The phrase "we carry it" cannot typically be broken up in this way without becoming ungrammatical or unacceptable. From my observations, co-text emotes can appear between any words or phrases without issue. They are not syntactic elements, instead iconically modifying their word or phrase. I posit that co-text emote usage is akin to the use of unusual fonts for humorous or satirical purposes.

(8) amazon prime 7/10/24, 8:38 PM
you'll want to look for trad wife `woman's duty` books

In the above example, stylizing the phrase "woman's duty" as in-line code visually separates it from the rest of the utterance and emphasizes the user's sarcasm. Removing that stylization fully preserves the meaning of the text, and the same can be said of removing co-text emotes. Below are versions of (6) and (8) in a typical style:

(9) For the love of God can't we love each other just a little?

(10) you'll want to look for trad wife woman's duty books

Each of these modified messages convey exactly the same meaning as the originals. They are perhaps less funny, because there is no absurdist or sarcastic visual cue embedded in the text. As we have established before, removing a pre- or post-text emote *does* affect the meaning of an utterance. Example (5) is restated for convenience and is then modified to remove the emotes:

(11)

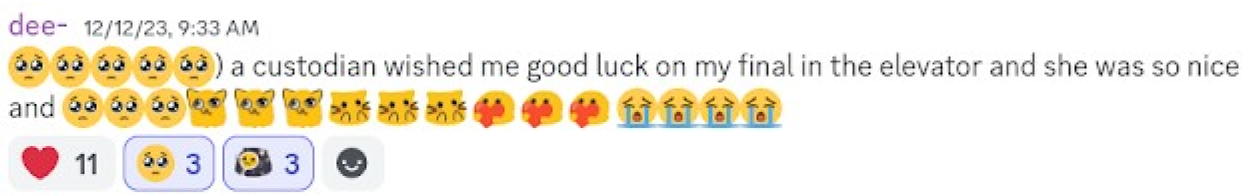

(12) a custodian wished me good luck on my final in the elevator and she was so nice and

Clearly, there is a significant difference between (11) and (12). The pre-text pleading emotes 🥺 serve as parallel emotion signals (vii), and the long string of post-text emotes serve as emotion signals (vi) and as emotion intensity enhancers (viii). The utterance describes in three parts an event, the quality of the event, and the emotional response of the utterer. Without the emotes, the utterance is missing the entire third part representing the utterer's emotional response. This information is not recoverable using what remains. Additionally, reducing the quantity of the pre- and post-text emotes also changes the meaning of the utterance: if there are fewer "intensity enhancer" elements, the utterance will be less intense.

This pattern of meaning can be applied to all pre-, post-, and co-text emotes: removing co-text emotes does not change the meaning of the utterance, but pre- and post-text emotes may not be removed without affecting the meaning. Co-text emotes can be better understood as a stylistic choice used to indicate humor, similar to choosing unusual fonts solely for the visual effect. Post-text emotes, on the other hand, are comparable to Harris' 1982 analysis of conjunctions. He names the operators "not", "and", "or", and the semicolon as syntactic "residues" operating on a sentence pair. While functioning to connect two utterances, "...[they] can appear not only after S_1 but also at interior points of S_1." The post-text *pleading* 🥺 illustrates this principle. It functions as an emotive conjunction between the two utterances. It cannot easily be analyzed as a "residue" since it does not readily correspond to a phrase such as "I deny" or "I disjoin". Storment (2024) does provide a formal logic representation of *pleading* 🥺 (p. 35), reproduced below:

(13) ∃e[making the 🥺 facial expression(e) & Agent(e, x) & Goal (e,y)]

The emote 🥺 in (11) cannot be confidently paired with a zeroed "I plead", since in this case the utterer is not pleading. They're expressing a more general sentiment of affection and being (emotionally) touched. Storment provides the formal logic evaluation in (13) because *pleading* 🥺 has no lexical equivalent, at least in English. This is true of many emotes, and is likely a contributing factor to the practice of orthographic and verbal representation (to be discussed elsewhere).

Harris (1982) defines the concept of a *daggered sentence*, which is a sentence reduced from a necessary but grammatically impossible form. One cannot say "Do you want tea I disjoin coffee?" in place of "Do you want tea or coffee?", but the former is necessary to interpret the "or" present in the latter. I propose there exist such daggered sentences for emotes. Consider the utterance "i can't believe you just did that 😱 😡". This is a "daggered" equivalent of the impossible utterance "i can't believe you just did that shock anger". Co-text emotes are not daggered equivalents of utterances since they do not contribute syntactic or semantic meaning.

Finally, there is a strict divide between the "pools" of emotes appropriate for use as co-text or as pre- or post-text. Co-text emotes can be essentially any universal emote. The majority of universal emotes are objects, and these are heavily favored in co-text usage. Co-text emotes usually depict objects, activities, or symbols. A select few universal emotes may be used in both contexts but are preferred as pre- and post-text emotes. These include (but are not limited to) hearts, sparkles ✨, skull 💀, eyes 👀, and facial expressions such as *pleading* 🥺, *sob* 😭, *flushed* 😳, and *sunglasses* 😎. I have not generally observed any custom emotes used in a co-text manner, but it may be possible (if rare). Post-text emotes function more as "tone indicators" and are usually facial expressions, characters, or other semantically meaningful icons. Object emotes do not appear in pre- or post-text settings. In addition, homophony and wordplay are never employed in pre- and post-text usage.

As an additional diagnostic, any emote found in pre- and post-text environments may also appear in pro-text environments. In the '24 server, acceptable pro-text emotes are almost exclusively custom emotes. Storment (2024) provides several examples of iconic or "object" universal emotes in pro-text usage, but conventions within the '24 server do not allow for iconic pro-text emotes. The same handful of universal emotes acceptable for pre- and post-text use are similarly acceptable as pro-text emotes, though possibly at a lower frequency than custom emotes. They are also represented orthographically in the same usage patterns as custom emotes.

Differences Between Pre-text and Post-text Emotes

Pre- and post-text emotes serve a function similar to sentential adverbs, and can be similarly analyzed. The placement of the emote either before or after the sentiment may correspond to certain roles of Li & Yang's (2018) framework, but further investigation is needed. These roles are restated below for convenience:

(i) attitude signal

(ii) attitude intensity enhancer

(iii) illocutionary force modifier

(iv) humor

(v) irony

(vi) emotion signal

(vii) parallel emotion signal

(viii) emotion intensity enhancer

Post-text positioning functions as a *tone indicator*, which are widely used in various forms. The most common forms do not employ emotes at all and are enclosed in parentheses or preceded by a forward slash.

(14) a. you could also go to the gym /hj
or play smash with me /srs

b. u all have science-brain (derogatory)

Forward-slash tone indicators have a handful of canonical forms, including those in (14a). These forms are often abbreviations of common words or phrases; /hj is "half-joking", and /srs is "serious". Parenthetical tone indicators are not generally abbreviated and have more freedom in what they can contain. When read aloud, a forward slash is realized as the word "slash". Parentheses are not realized in speech. While both of these forms have some degree of flexibility in what they can include, emotes as tone indicators have much more freedom. They do not have to represent any specific words or phrases, and may give the utterance an emotional impression that would be difficult to express in words.

Another form of tone indicator is *bracketing*, or enclosing an expression with symbols. This is distinct from a parenthetical expression and is more akin to italicizing or bolding text. Bracketing with double forward slashes, as in (15a), is not a forward-slash tone indicator. (15b) utilizes Markdown to create in-line code formatting. Emotes as brackets can be universal or custom. Bracketing emotes have a different function from post-text emotes and can be applied to only a portion of an utterance.

(15) a. //ok myrin//

b. but the bottom ones `were` impacted

c. ✨homosexual tendencies✨

d. where im from will become `where do i live` vs `where am i *from*`

Multiple styles of bracketing can be combined for further nuance, as in (15d).

Self-reacts

A subset of type (iii) emotes modifying illocutionary force are *self-reacts*, or when the sender of a message "reacts" to their own message. On Discord and other platforms, a user may "tag" a message by attaching one or more emotes. These emotes appear below the message in boxes, with a number signifying how many users have reacted with that emote.

(16) 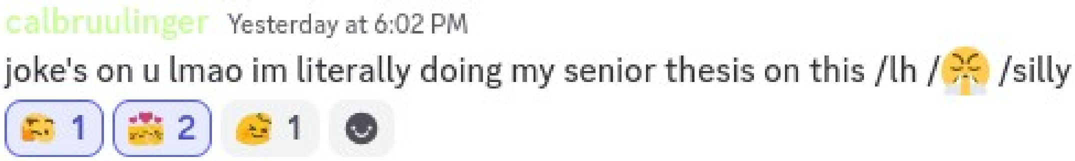

In the above example, the rectangles beneath the message are *reacts*. These can be universal or custom emotes (all of the reacts in the above example happen to be custom emotes). The sender of this message *self-reacted* with *thnark* and *meowhuggies*. *Thnark* is an expression of disapproval, questioning the reasoning behind the utterance it operates on. *Meowhuggies* is a hug, or more generally an expression of empathy and comfort. Combining the two suggests a compound interpretation of "thnarkhuggies", conveying both disapproval and a hug. Another user also reacted with *meowhuggies*, indicating their agreement with the sentiment of *meowhuggies* and making the count 2. That user also reacted with *headpat* (which indicates the reacter patting the sender on the head).

Self-reacts are an important aspect of hedging one's utterances, by indicating the intended attitude of the message. Most self-reacts are negative in connotation, inviting readers to take the message unseriously while acknowledging the unfortunate content of the message itself. *Self-thnarking* is by far the most common self-react, in which a sender reacts to their own message with *thnark*. This "invit[es] others to react sardonically and with mild disappointment or discontent at [the sender's] own messages".[2] The sender in example (16) wanted to protect against a passive-aggressive interpretation of their message, employing slash-style tone indicators and self-reacts. This message was in reply to a different user's message from spring of

2021, and a reply years later is likely to be taken poorly unless extreme hedging is present.

The first attested usage of *self-thnarking* is from late summer of 2020, very early in the server's lifetime. The user wants to communicate that he is making bad suggestions, and that he is aware of this. In this case, the self-thnarks fall into types (iii) and (v), type (v) being an indicator of irony.

(17)

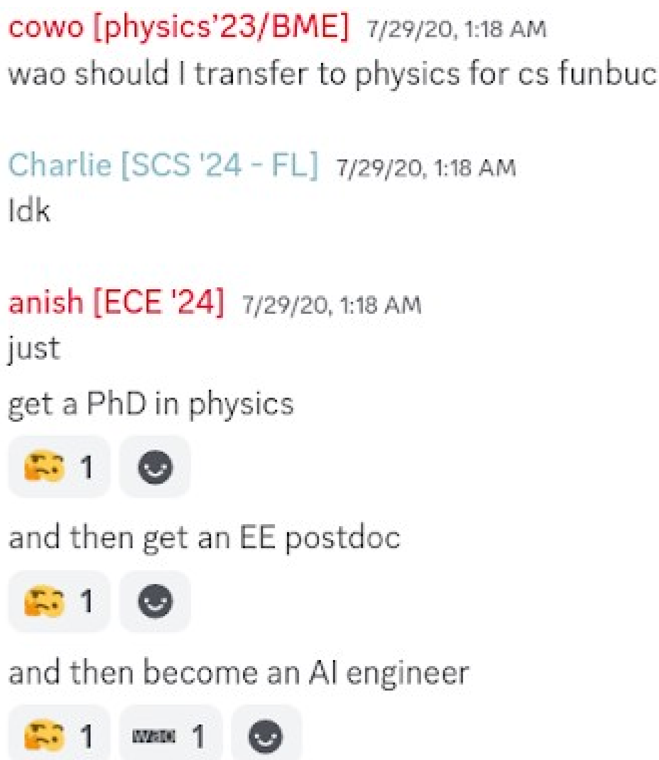

The conversation then turned to the self-thnarks themselves, and this continued to be a topic of interest and discussion in the server. Users often sent messages such as "don't self-thnark", indicating that they thought the sender's message was perfectly acceptable. Self-thnarking soon became an emote of its own, as the image of the reaction box containing the emote. This was mostly created for amusement and treated as unserious, but another layer of recursion was discussed (and implemented):

(18) 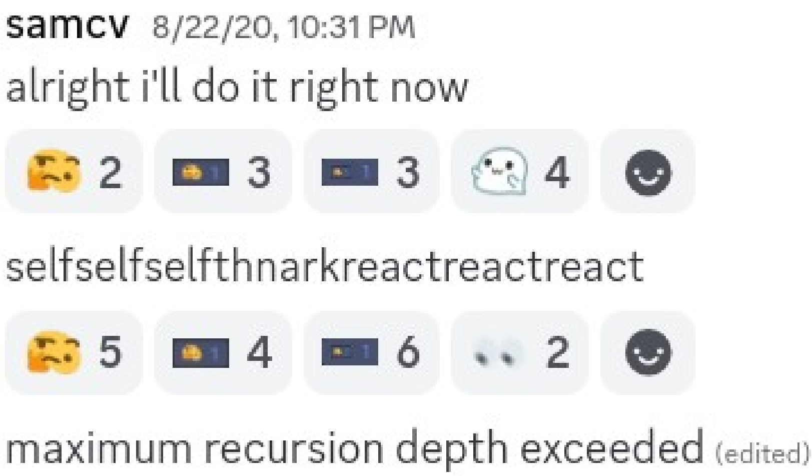

The react to the immediate left of the *thnarks*  is a "self-thnark" emote (19), which is a screencap of what the react looks like. To the left of *self-thnark* is "selfselfthnarkthnark" (20), a screencap of a *self-thnark* emote react.

(19) 

(20) 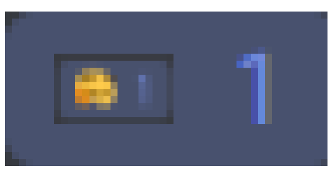

Due to the small maximum filesize allowed for emotes, it is difficult to discern the features of (19) in (20) and the image is of low quality.

As discussion about the phenomenon evolved, users began to type the phrase "self thnark" in the body text of their messages. This did not usurp self-reacting, but it allowed for a generic "self-[ ]" construction into which anything could be placed. Users also took advantage of the mechanic as a vehicle of humor, type (iv), without falling into types (iii) or (v).

(21) 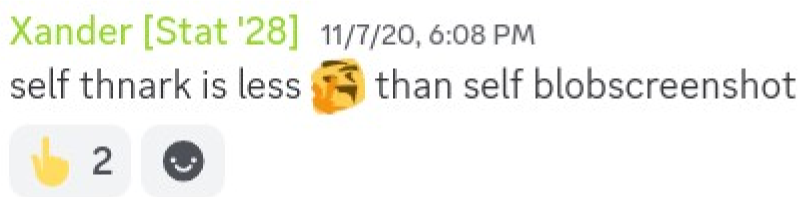

(22) 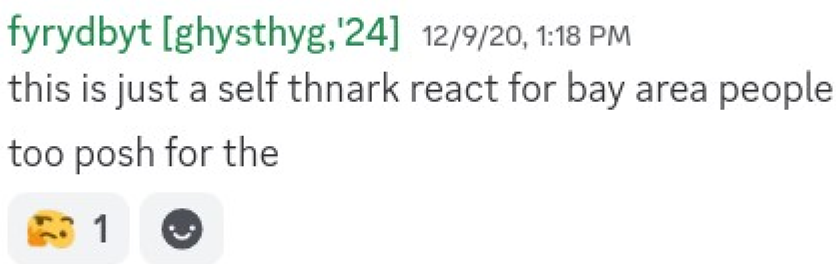

The message in (21) contains *thunj* as a pro-text emote. In this context, *thunj* has a sense of "inappropriate" or "unacceptable" behavior. (the reacts present are not self-reacts). In (22), the sender is not self-thnarking to indicate irony or to alter the illocutionary force. They are "finishing the sentence" started in the second line of the message, intending it to be read as "too posh for the self-thnark". This is a rare occurrence when compared with other self-reacts, as most fall into types (iii) and/or (v).

Users never explicitly defined the meaning of specific self-reacts, but there is a consensus on the general connotations of each. Broadly speaking, a self-react can be interpreted literally as reacting to one's own message, under the pretext that the message is sent by someone else. Self-thnarking in this case is the sender reacting with disapproval or disappointment at their own sentiment. This can be extended to other custom emotes with complex meanings, and can add nuance to one's tone in a way different from a tone indicator.

(23) 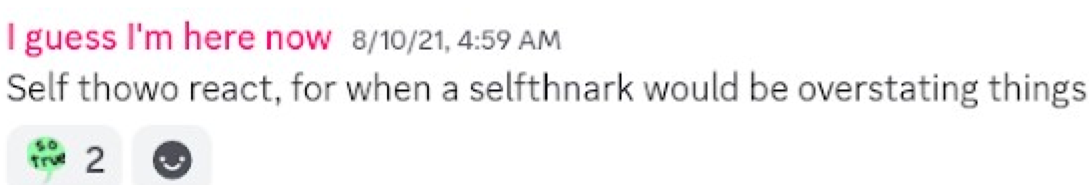

This user does not explicitly define the meaning of "self-thowo", but suggests that it is a type (iii) modifier weaker than self-thnark. I would interpret a "self-thowo" to mean that the sender is intrigued by their message, subtly inviting others to engage with them on the subject.

When used as a sentential emote, *thunj*  can be interpreted as a more aggressive version of *thnark* . Not only does it express disapproval, it often expresses contempt. Appearing in both pre- and post-text environments, the placement of *thunj*  determines the recipient of contempt. Pre-text is usually directed at another user, and the sender's message is a response to something that user has sent. Post-text is a more concrete example as outlined in Storment (2024), and is intended to modify the tone of the sender's message. In this case, post-text emotes are type (viii) and serve to intensify emotion in the message. Pre-text emotes are harder to categorize, but would perhaps fall into type (vi), serving as emotion signals.

Differences Between Tone Indicators and Self-reacts

Since tone indicators and self-reacts are both used to hedge an utterance, one might argue that they shouldn't be considered distinct mechanisms. This argument can be easily disproven, since tone indicators and self-reacts have different patterns of use. If they served the same function, they would appear interchangeably and without pattern. It is true that both mechanisms perform similar functions. The difference lies in levels of scope. Tone indicators hedge the content of an utterance, while self-reacts hedge the act of utterance.

This separation occurs frequently in natural language. For example, including the word "honestly" in one's utterance is part of that utterance's content. Actually telling the truth is part of the act of utterance. The use of the word "honestly" does not obligate the utterer to be honest. You can say "I'm telling the truth" and lie, and you can say "I'm lying" and then tell the truth. These additions are tools to hedge the content of an utterance. Hedging the act of utterance in natural language often involves co-utterance gesture. If an utterer makes a statement that they want their addressee(s) to agree with, they may nod their head while making that statement – this influences the act of utterance but is not part of the content.

Sent messages in Discord may be edited by the sender at any time, and elements such as tone indicators can be added, changed, or removed. These alterations are post-utterance elements affecting the content of the utterance. Users can also react to sent messages at any time (their own and others').

Another important difference between tone indicators and self-reacts is exemplified by the sender's hypothetical capability to perform the act on another person's message. Reacting to others' messages is similar to reacting in live-language situations. If someone says something I agree with, I can nod my head to indicate my agreement. If someone sends something I agree with, I can react with a thumbs-up or similar positive emote. No matter the medium of communication, these "other-reacts" are perlocutionary effects. It follows that self-reacts are emulations of perlocutionary effects. In contrast, tone indicators may only be produced by the utterer. I cannot influence the tone of someone else's utterance online or offline. In employing a self-react, I am indicating how I want others to respond to my utterance. Tone indicators are used to affect others' interpretation of my utterance.

Syntactic Analysis

The syntactic status of pre- and post-text emotes is made clear once they are included in syntax trees. Categories strictly follow a hierarchical organization, with self-reacts having the broadest scope, and pre-text emotes the narrowest:

self-react > tone indicator > post-text > pre-text

Each category is represented by a node which moves at LF to take appropriate scope. This process is well-attested in natural language – consider the negation in "I'm <u>not</u> going to eat lunch" and in "I'm going to <u>not</u> eat lunch". The same unit, "not", yields different interpretations depending on its scope. This is also true for emotes. They are indivisible units that perform distinct functions at each 'level' of scope.

Pre-text emotes act in a similar manner to interjections preceding an utterance (such as "hey" in English). They remain at the beginning of the sentence and occupy a sibling node, as below:

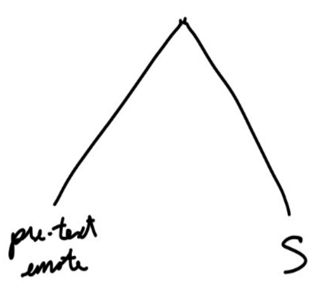

Post-text emotes must be moved from the end of the sentence to the front of the tree in order to gain the semantic scope over pre-text emotes:

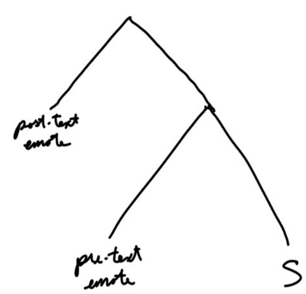

Tone indicators, regardless of style, must also move to the front of the tree. If there is more than one tone indicator, they move as a unit containing all tone indicators.

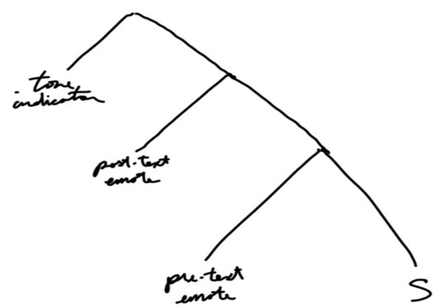

Self-reacts command *all* of a message's content, since they are a remark on the utterance as a whole. I propose a new category within which to include self reacts – *post-utterance* elements. Below is a tree with all four elements (this is possible, but I have not yet observed such an utterance).

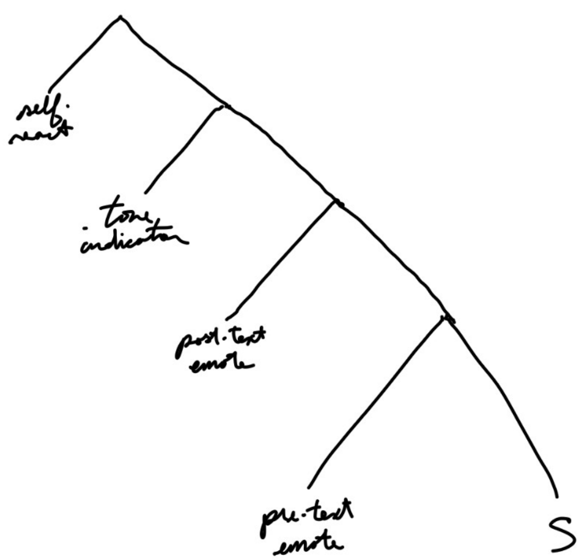

Let us apply this analysis to exemplify its mechanisms. Below are two complex utterances employing multiple categories of emote use. (24) contains both a pre-text and a post-text emote. (16) is repeated here as (25) for convenience, since it contains tone indicators and self-reacts (the in-line emote is part of a tone indicator and is not a post-text emote).

(24) 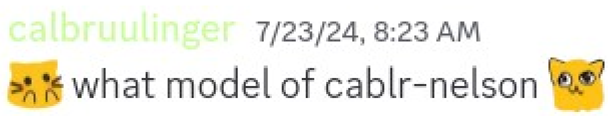

(25) 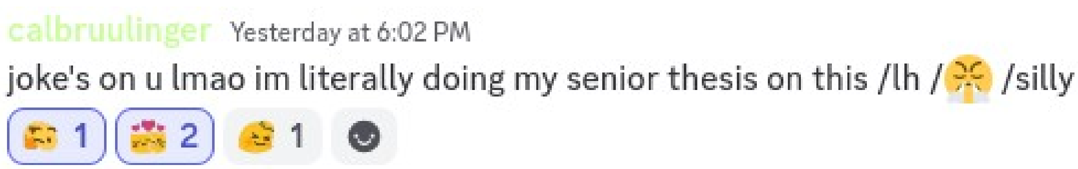

In example (24), the pre-text emote is a "part" of the sentence on the same level as the text. The post-text emote applies to the whole sentence preceding it, *including* the pre-text emote. A reasonable tree might be drawn as follows:

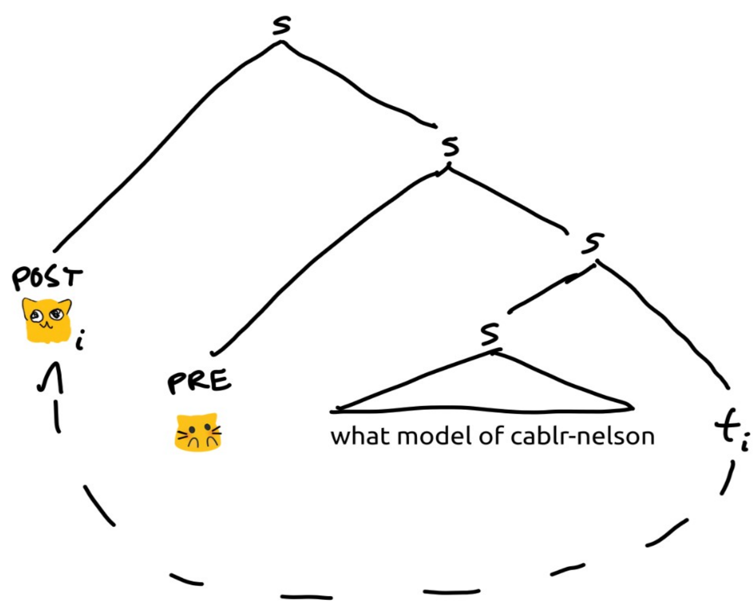

The pre-text and S nodes are siblings, and the pre-text emote c-commands the text portion of the sentence. The post-text emote shifts from a sentence-final position to the front of the tree, having scope over both the pre-text emote and the text portion.

Example (25) is more complicated, since it has multiple tone indicators and multiple self-reacts. The tone indicators shift as a unit, and each element remains in the same order as it appeared in the utterance. The self-reacts cannot shift in the usual sense since they are not part of the body of the message. They exist outside of the content and are applied to everything within the message, indicating the utterer's attitude toward their own message. A tree for this example is below:

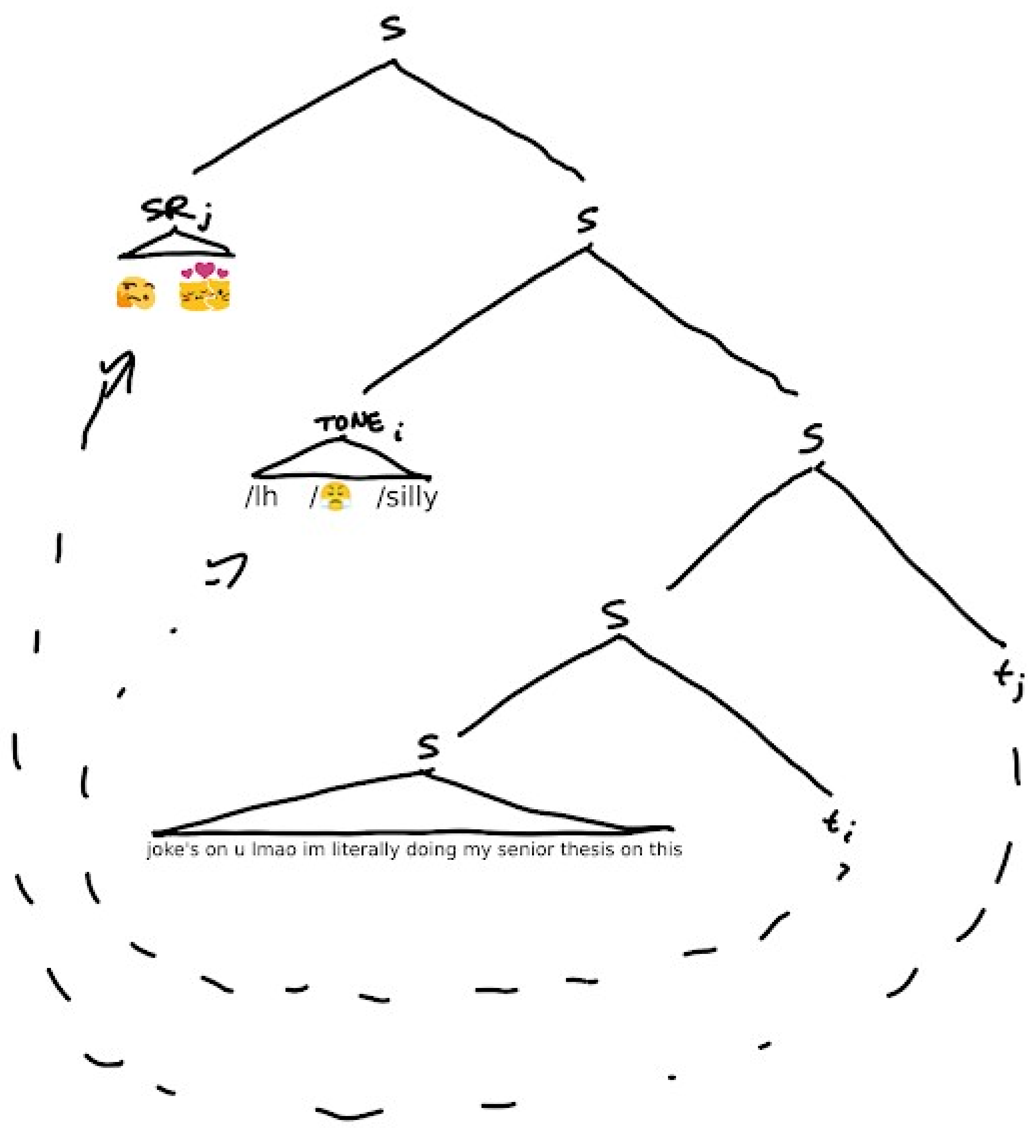

Conclusion

In this paper I have classified emote use into four types, each with different syntactic properties and semantic functions. I have introduced *act-of-utterance modifiers*, a class usually not possible in live language in person. This is an area ripe for further exploration. Division of such a class into perlocutionary *effects* and perlocutionary *acts* could prove a valuable framework. Analysis of more general emote usage as linguistic signalling is also a potential area for future work, especially concerning complex or compound emotes such as the recursive self-thnarks.

APPENDIX

Glosses of select examples

(3a) [positively emotionally touched/affected, in a frolicsome way] i think i might really like to……………catch up on schoolwork waaah [humorous despair] once i'm back on top of stuff though? [appreciative][contented]

(4a) taiwan is so nice <u>by the way</u> [appreciation x4] i like it here <u>so much</u>

(5/11a) [positively emotionally touched/affected x5] a custodian wished me……………she was so nice and [tearing up from being touched/"that's so sweet" x3][frolicsome pos. emotionally affected x3][pos. soft gasp x3][appreciative x3][happy cry x4]

(16/25a) joke's on <u>you</u> [laughing my ass off] i'm literally………
[lighthearted][triumphant/"hell yeah"][silly]

self-thnark → disapproval

self-meowhuggies → hug (combines with self-thnark)

(18) alright i'll (<u>make the recursive emotes</u>) right now
self(self(self-thnark react)react)react

<u>1st level:</u> self(self-thnark react)react → self-react of emote: image of self-thnark

react

<u>2nd level:</u> self(self(self-thnark react)react)react → self react of emote: image of *image of* self-thnark react (image of 1st level)

(21) self thnark is less [disapproval + contempt] than [inviting others to save/take special note of message]

(24) [positive soft gasp] what model of cable-nelson (<u>piano</u>) [soft positive "oh my gosh"]