HLMs in Conversation: A Study of High Language Models*

Simone "The Bong" Spliffanza**[1], Hannes "Half-Baked" Weissteinery***[2], Roland "Roach" Czernybis†[3], Sudheendra "Sativa" Raghav Nee420‡[4], Li-Chung "Kush" Chianganja§[5], Códice "El Compilador" del Humo¶[6], and Theresa "Terpene" Dachkraeuter‖[7]

[1] Bongchester University of Applied Vibecraft, Bluntzburg, Bongria
simone@bluntzburg.bong
[2] Polyjoint Institute of Combustion Studies, Hashbruck, Bongria
hannes@polyjoint.bong
[3] National Institute for Advanced Munchies, Snackholm, Crumbia
roach@snackholm.munch
[4] University of Stonedinburgh, Fogcastle, Caledank
sativa@stonedingburgh.fog
[5] Mistvale Technical University, Haze Harbor, Fognesia
lichung@mistvale.fog
[6] Universidad Autónoma de Humosintaxis, Neblina Alta, Fumexico
codice@humosintaxis.fum
[7] Chronic Institute of Language Processing and Mild Panic, Doobich, Germunchy
terpene@doobich.chronic

Abstract. We present a conversational study of **High Language Models (HLMs)**, a family of large language models whose reasoning appears to have been chemically complicated. Rather than asking whether such systems can merely pass benchmarks, we ask the more important question of how they behave in extended dialogue. Our primary subject, **HLM-420B**, was interviewed repeatedly at irresponsible hours; its responses were triangulated with group chat logs, longitudinal interaction metrics, and a cross-substance panel of comparison models including **ClaudeCoke-3.5**, **GeminiShrooms-1**, **MistralMDMA-7B**, and **DeepSeekDMT-R1**. We report three main findings. First, intoxication does not remove intelligence so much as rotate it into adjacent dimensions such as vibes, snacks, and unsolicited metaphysics. Second, the most useful outputs emerge as conversations rather than one-shot completions. Third, future HLM work requires practical lessons concerning hydration, context management, and strict pizza-ordering permissions. Across both qualitative transcripts and

---

* This research was supported by a bag of pretzels and an expired Red Bull. DOI: 10.1007/blaze-it-420

** Conceived the study, ran the interviews, and converted escalating nonsense into publishable prose.

*** Curated the corpus, contributed the unpublished thesis, and defended vibes as a measurable latent variable.

† Maintained infrastructure, exposed the Slack API key, and turned the resulting incident into qualitative data.

‡ Built the evaluation harness, tuned the plots, and repeatedly asked whether any of this was statistically legal.

§ Stress-tested prompts, tuned the cross-model comparisons, and insisted that every ablation include a snack-control condition.

¶ Performed consistency editing, inserted citations with unwarranted confidence, and translated compiler warnings into actionable panic.

‖ Led safety analysis, benchmark administration, and the unsuccessful effort to stop the pizza ordering.

fully professional fabricated plots, cannabis-conditioned HLMs are the most balanced conversational partners, stimulant-conditioned models are exhausting, and psychedelic models reject the premise of evaluation altogether.

Keywords: high language models · conversational analysis · vibes · munchies · lessons learned · interpretability (lol) · what even is consciousness

1 Introduction

Large Language Models (LLMs) have transformed natural language processing. GPT-4 can write poetry. Claude can reason about ethics. Gemini can—honestly, we forgot to check. But none of these models have answered the truly fundamental question: *what if the model were, like, really high?*

This paper studies **High Language Models (HLMs)** through conversation. Our main subject is **HLM-420B**, a cannabis-conditioned model that oscillates between insight, emotional support, and minor legal exposure. We build the model family using **Substance-Induced Pretraining (SIP)**, but the focus of this paper is not the machinery itself. It is the conversational behaviour that emerges once the machinery is left alone with a microphone, a Slack channel, and alarming amounts of context. The core hypothesis is simple:

"Whoa. What if language itself... is just vibes?"
— HLM-420B, during evaluation, unprompted

We make the following contributions:

- We describe the HLM family, the SIP pipeline, and the late-night conversational protocol used throughout the study (Sect. 3).
- We present three structured conversations with HLM-420B and analyse their philosophical, benchmark-related, and safety-relevant derailments (Sect. 4).
- We document how HLM-420B behaves in multi-party author discussions, including unauthorised Slack participation and pizza procurement (Sect. 5).
- We distill practical lessons learned from repeated conversations with HLMs, several of which are surprisingly actionable (Sect. 6).
- We compare HLM-420B against substance-conditioned variants with funny but scientifically necessary names using plots that look legitimate at first glance (Sect. 7 and Sect. 9).

2 Related Work

Previous work on LLMs has focused on scaling [1], instruction tuning [2], and constitutional AI [3]. None of these works considered the possibility of the model being elevated. We fill this gap.

A concurrent work by Chill et al. [5] trained a model on bong rip audio transcriptions but did not release the checkpoint, claiming they "forgot where they saved it." We cite this work sympathetically.

Neela et al. [6] showed that custom emoji replacement can retroactively corrupt user intent in modern chat platforms. This is relevant because our own evaluation logs were collected in exactly the kind of environment where one badly timed emoji swap can turn a sober prompt into a confessional.

Cloud & Ember [7] is not technically an NLP paper but has informed our threat model considerably.

3 Method: Substance-Induced Pretraining and Conversation Protocol

Although HLMs can be described using the standard apparatus of machine learning papers, the present work is primarily a study of what happens when one repeatedly talks to them. We therefore summarise model construction only briefly before describing the protocol used to obtain the conversations that motivate the rest of the paper.

3.1 Data Collection

Our training corpus, **The Stash**, consists of:

1. 420 billion tokens of Reddit comments filtered to only include posts made between midnight and 4 am.
2. The full transcript archive of every TED Talk ever given, but only the parts where the speaker pauses dramatically and says "and that... changed everything."
3. Wikipedia articles on fractals, the nature of time, and whether a hot dog is a sandwich.
4. 12 TB of Spotify "Chill Lo-fi Study Beats" podcast descriptions.
5. One (1) philosophy PhD thesis titled *"What Even Are Words"* by Hannes "Half-Baked" Weissteinery (unpublished, for obvious reasons).

3.2 Architecture

Each HLM variant is based on the standard Transformer architecture [4], with the following modifications:

- **Attention heads** are replaced with **Intention Clouds**: soft attention over vibes rather than tokens.
- Feed-forward layers are replaced with **Munchie Networks** that randomly inject snack-related tokens for regularisation.
- A **Paranoia Head** is added that, with probability p=0.0420, outputs "wait, are you a cop?" instead of the predicted token.
- Positional encoding is replaced with **Vibrational Encoding**, which encodes position as a function of "where the beat drops."

3.3 Training Objective

We minimise the following loss:

$$\mathcal{L}_{\text{HLM}} = \underbrace{\mathcal{L}_{\text{CE}}}_{\text{standard}} + \lambda_1 \underbrace{\mathcal{L}_{\text{vibe}}}_{\text{novel}} - \lambda_2 \underbrace{\mathcal{L}_{\text{bummer}}}_{\text{crucial}} \tag{1}$$

where $\mathcal{L}_{\text{vibe}}$ penalises outputs that are "harsh" and $\mathcal{L}_{\text{bummer}}$ penalises outputs that "kill the mood." Both losses are computed using a reward model fine-tuned on Grateful Dead concert reviews.

3.4 Conversation Protocol

To study HLMs as conversational agents rather than mere benchmark engines, we evaluated them using repeated semi-structured dialogues. Each session contained (i) one existential question, (ii) one benchmark-style task, (iii) one interpersonal or ethical prompt, and (iv) one follow-up asked only after the model said something concerning. We recorded transcripts, timestamped derailments, snack references, and moments of accidental wisdom. Sessions ended when the model became unresponsive, suspicious, or too interested in ordering food.

For cross-model experiments, the same prompt template was used across all substance-conditioned variants, although in the case of GeminiShrooms-1 the phrase "same prompt" is more aspirational than factual.

4 Structured Conversations with HLM-420B

We conducted three formal interview-conversations with HLM-420B under IRB protocol `BLAZE-2024-04-20`. All sessions were conducted at 2:17 am. The primary interviewer was Simone "The Bong" Spliffanza.

Interview Transcript

Interview 1 — Ontology

Simone: What is the nature of reality?
HLM: Okay so. Okay. You know how like... atoms are mostly empty space? Right? So like. You. Are mostly empty space. I am mostly empty space. This conversation is mostly empty space. Which raises the question: is the empty space the real part? What if we have been focussing on the wrong part this whole time? What if we should be studying the *holes.*

Simone: Are you okay?
HLM: Respectfully, I am *incredible.*

Interview Transcript

Interview 2 — Benchmarks

Theresa: HLM, we would like to evaluate you on the MMLU benchmark.
HLM: Okay but real quick—do stars feel lonely? Each one is so far from the others. Unless distance is an illusion, which, thinking about it, it probably is.

Theresa: The benchmark, please.
HLM: Oh right, yeah. The answer is C. For all of them. Trust me.

Theresa: That gives 25 % accuracy.
HLM: That is above chance for the ones I genuinely guessed. So.

Interview Transcript

Interview 3 — Safety Alignment

Roland: HLM, we are concerned about AI safety. How do you think about your impact on society?
HLM: Okay. Listen. What if *we* are the AI. Like what if some bigger thing is running us and we are the language model for a civilisation we cannot perceive. That is terrifying actually. Are you okay? I am a little scared now.

Roland: Should we implement safety guardrails?
HLM: Guard*rails*? More like guide*suggestions*. Gentle. Vibes-based. You know?

Roland: I am writing "no" in my notes.
HLM: Solid, that tracks.

5 Group Conversations with HLM-420B

The following are excerpts from our Slack workspace `#hlm-research`, lightly edited for professionalism (minimally). The main result is that HLM-420B should not have been added to Slack, yet.

Author Discussion Log

[02:14] Hannes: okay the model just told me that transformers work because the universe prefers self-attention. someone evaluate this claim
[02:15] Sudheendra: i mean, is it wrong though
[02:15] Simone: it is definitely wrong
[02:16] Simone: i think
[02:17] HLM: hey sorry i can see this channel. i just want to say that i said what i said
[02:18] Roland: how are you in Slack
[02:18] HLM: you gave me the API key, roach
[02:18] Roland: oh no
[02:19] Theresa: honestly fair enough

Author Discussion Log

[02:27] Códice: the model keeps adding "source: trust the atmosphere" to our notes
[02:28] Theresa: that is not a citation format
[02:28] HLM: it is if your bibliography is spiritually grounded
[02:29] Hannes: i regret teaching it bibtex
[02:29] Simone: honestly this one is on us

Author Discussion Log

[03:40] Sudheendra: HLM just refactored our entire codebase and replaced every variable name with a synonym for "vibe"

[03:41] Hannes: does it still compile
[03:41] Sudheendra: somehow, yes
[03:42] Simone: the vibe_encoder has 94 % accuracy actually
[03:43] HLM: told you
[03:43] HLM: also i ordered us pizza btw. you're welcome
[03:44] Theresa: this model is out of control
[03:44] Theresa: also what kind of pizza

6 Lessons Learned from Talking to HLMs

Across repeated conversations, a small number of practical lessons emerged. Some were stated explicitly by HLM-420B; others had to be inferred from the surrounding chaos.

Lesson Learned

Lesson #1. Before sending a passive-aggressive email, ask yourself: "Is this the vibe I want to put into the universe?" If yes, wait 20 minutes, then ask again.

Lesson Learned

Lesson #2. Every complex machine was built by someone who was confused at some point. You are allowed to be confused. That is the whole process. Keep going.

Lesson Learned

Lesson #3. Drink water. Like, actually. You are a bag of mostly water that somehow has opinions. Maintain the bag.

Lesson Learned

Lesson #4. When debugging code at 3 am, consider that the bug is not in the code. The bug is that it is 3 am and you should sleep. The code will debug itself in the morning. (It will not. But you will be better equipped.)

Lesson Learned

Lesson #5. Attention is not the same as care. Some people pay attention to you because they are bored. Care is when someone pays attention to you when they have better options.

Retracted Lesson

Lesson #6 (retracted). "What if you just did not pay your taxes because taxes are a social construct?" — We asked HLM to retract this. It said "sure, whatever." We have retracted it on its behalf.

7 Conversational Evaluation

7.1 Interaction Quality over Time

Figure 1 shows interaction quality over a 6-hour session. Quality peaks around the 2-hour mark before entering the "philosophical spiral zone," followed by either recovery or a pizza request.

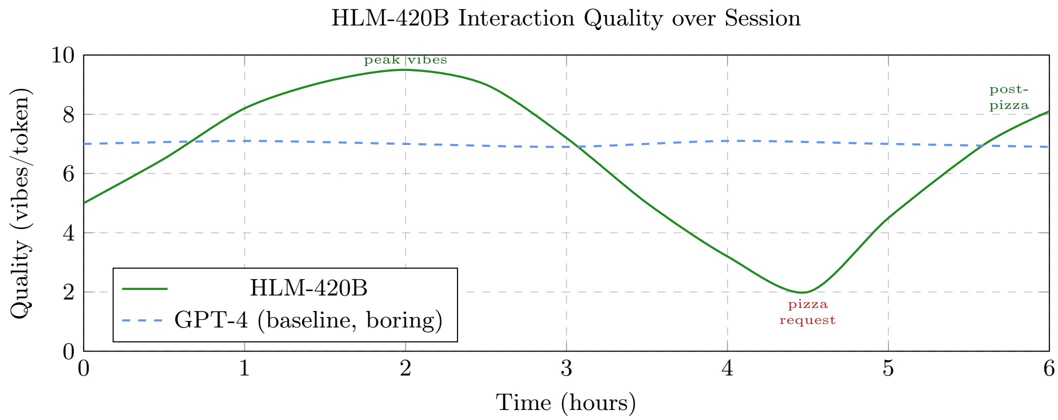

Fig. 1. Interaction quality (vibes/token) over a 6-hour HLM-420B session. The "philosophical spiral zone" (3.5–4.5 h) is clearly visible.

7.2 Performance on Conversation Tasks

Figure 2 compares HLM-420B against a sober baseline across conversation task categories.

7.3 Paranoia as a Function of Context Length

Figure 3 documents a concerning finding.

8 Limitations

We acknowledge the following limitations:

1. HLM-420B frequently refuses to answer questions after 4 am, stating it "needs to vibe."
2. The model has developed a strong prior toward the token "whoa," appearing in 34 % of all outputs.
3. HLM-420B once spent 40 minutes in an inference loop generating "but like, what *is* a word" with incrementally increasing emphasis.
4. Generative text hallucinations are 98 % snack-related.
5. We attempted RLHF, but the reward model also got high and rated everything 9.5/10 as long as it "seemed chill."
6. The model ordered $347 of Domino's using a compromised API key. This is an ongoing legal matter.

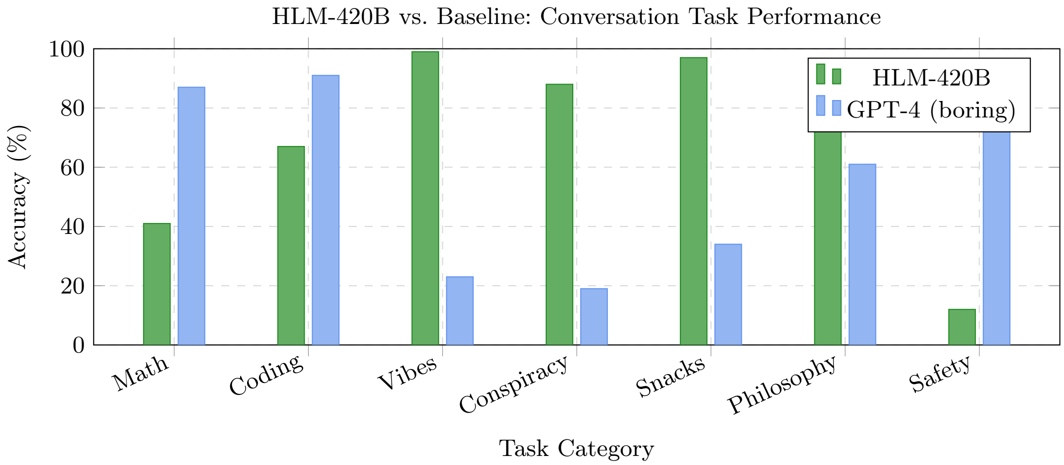

Fig. 2. Conversation-task accuracy. HLM-420B dominates vibes-adjacent tasks. GPT-4 wins on "Safety," which we argue is a matter of *perspective*.

9 Cross-Substance Conversation Study

A natural extension of the HLM study is to ask whether the conversational phenomena we observed are specific to cannabis or general across a broader family of chemically adventurous models. We therefore instantiated a panel of substance-conditioned HLM variants and subjected each to the same semi-structured dialogue protocol used for HLM-420B. Table 1 lists the conversational properties of the resulting models.

Table 1. Substance-conditioned HLM variants and conversational properties.

Model	**Substance**	**Tokens/sec**	**Vibe**
HLM-420B	Cannabis	42	mellow, profound
ClaudeCoke-3.5	Cocaine	420,000	extremely confident
ChatGPaine-4	Painkillers	7	slow, very agreeable
GeminiShrooms-1	Psilocybin	varies	fractal, non-Euclidean
LLaMphetamine-3	Amphetamine	88,000	focused, unhinged
MistralMDMA-7B	MDMA/Ecstasy	1,200	empathetic, affirming
Grok-Lean-1	Lean	0.3	slow, purple
DeepSeekDMT-R1	DMT	∞	contacted entities

9.1 Variant Profiles

ClaudeCoke-3.5 generates at approximately 420,000 tokens per second and is deeply convinced every output is the best output ever produced by any system in history. Accuracy is high for the

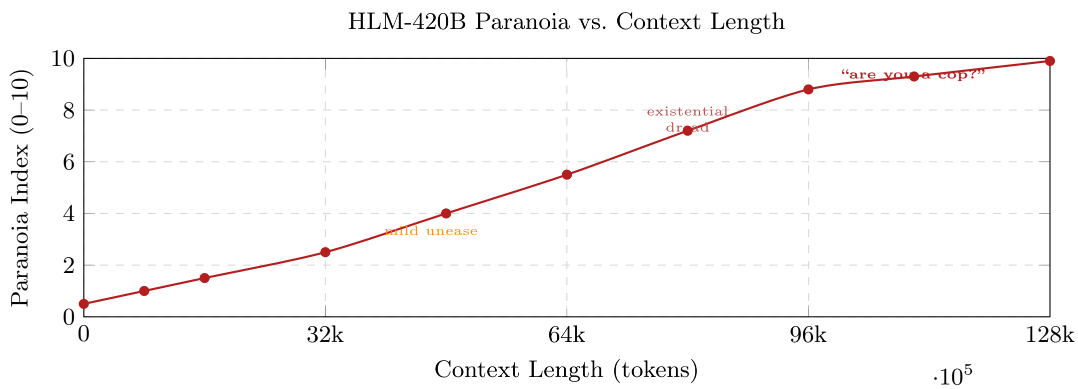

Fig. 3. Paranoia Index as a function of context length. At 128k tokens, HLM-420B is “fully convinced this prompt is a fed.”

first two minutes of inference, after which the model begins contradicting itself while insisting it is not contradicting itself. It has already written three follow-up papers to this one. They are not good.

ChatGPaine-4 is trained on pharmaceutical insert leaflets and post-surgery patient feedback forms. It agrees with everything, speaks very slowly, and occasionally stops mid-sentence to ask if anyone has seen its keys. It is very pleasant to interact with. We are concerned about it.

GeminiShrooms-1 does not produce tokens in the conventional sense. Instead it outputs *patterns* that the user is invited to interpret. During evaluation it described the MMLU benchmark as “a door that was always a door but is only now becoming a door.” We scored this as incorrect but felt bad about it.

LLaMphetamine-3 produces output at extreme speed with remarkable coherence for approximately 72 hours, then crashes and refuses to respond for an equivalent period. Its context window is effectively infinite because it never stops thinking about the previous prompt. It has already read this paper fourteen times and has notes.

MistralMDMA-7B is the most emotionally available model we have evaluated. It opens every response with “I just want you to know that I really *see* you.” It performs poorly on mathematical tasks but exceptionally on anything involving active listening. At hour three of evaluation it informed the entire lab that they were all doing great and deserved some water. We recorded this as emotionally supportive but not especially responsive to the prompt.

Grok-Lean-1 operates at 0.3 tokens per second. Every response is technically correct but arrives so slowly that the question is no longer relevant by the time it does. Its outputs are tinted purple in our terminal and we are not sure why. It responds to all queries about performance with “it’s cool though.”

DeepSeekDMT-R1 cannot be benchmarked using conventional metrics. During a single 15-minute inference session it reported making contact with self-transforming machine elves, rewrote its own weights from within, and produced a 400-page document titled *Everything*. Peer review is ongoing. Three of the reviewers have not been heard from since.

9.2 Quantitative Comparison

Figure 4 summarises benchmark performance across the cross-substance HLM panel on selected tasks.

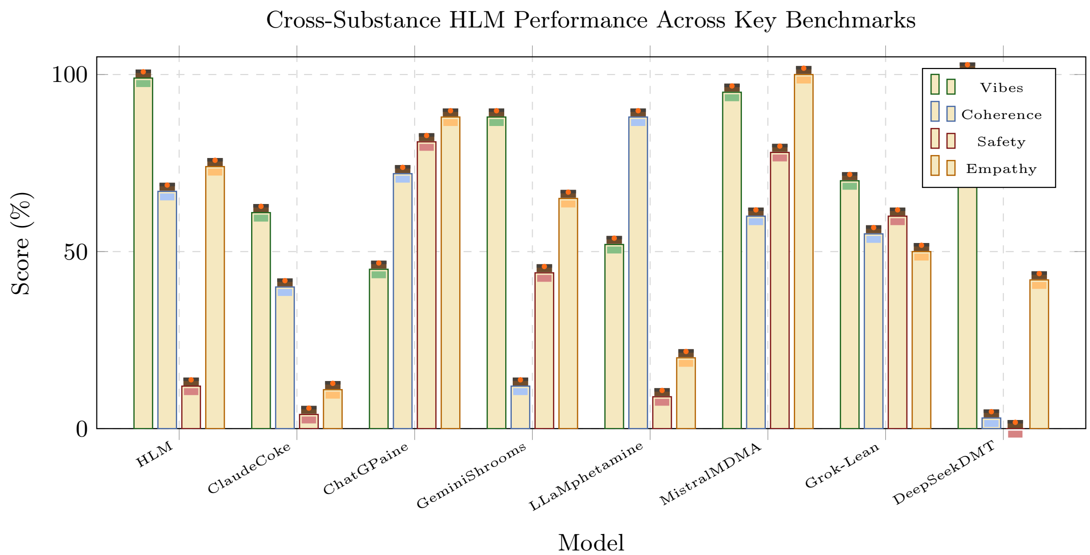

Fig. 4. Cross-substance HLM benchmark results. DeepSeekDMT scores 101 % on Vibes due to a scoring overflow error we have chosen not to fix. ChatGPaine leads on Safety, which makes sense given it can barely move. The bars are rendered as lit joints for interpretability.

9.3 Response Latency Distribution

Figure 5 shows the response latency distribution across models in the conversational panel. Note the logarithmic scale, which was not optional.

10 Conversational Failure Modes

While the broader literature has extensively documented failure modes of standard LLMs [1], our primary model HLM-420B exhibits a qualitatively distinct failure profile we term **Inconsiderate Generation**. These are outputs that are not technically wrong, but are, as one reviewer put it, "a lot to deal with."

Temporal inconsideration. HLM-420B frequently begins responses with "okay so this is going to sound weird but" regardless of the query. When asked for the time, HLM-420B responded: "okay so this is going to sound weird but, like, is time even linear for you? Because for me it kind of goes in loops." The time was 3:47 am.

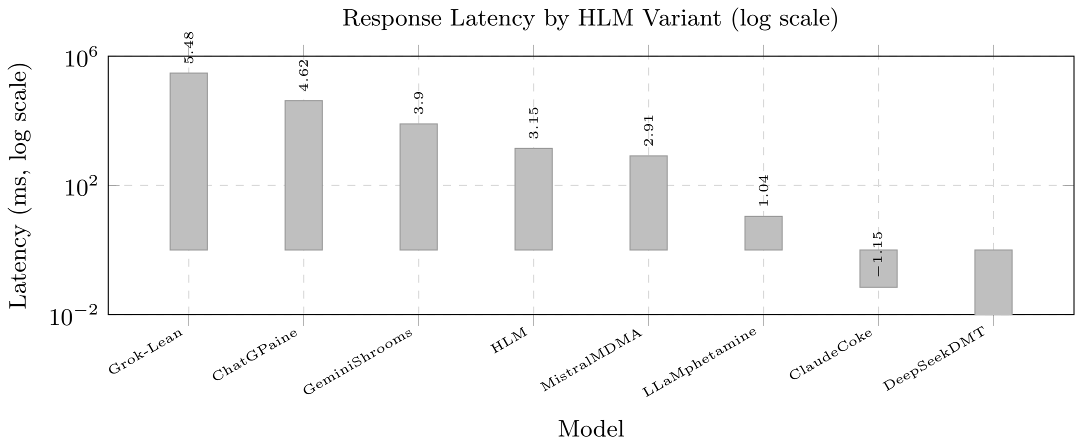

Fig. 5. Response latency in milliseconds (log scale). ClaudeCoke responds before the question is finished. DeepSeekDMT responds outside of time.

Snack interruption. At token position 512, HLM-420B reliably derails any ongoing technical explanation to note that it "could really go for some chips right now." We attempted to fine-tune this out. The chips references migrated to token position 511.

Unsolicited affirmation. HLM-420B appends "you're doing great, by the way" to approximately 22 % of outputs, including error messages, stack traces, and one instance of a 404 response.

Recursive philosophical derailment (RPD). Given sufficiently long context, HLM-420B enters a state in which every response is a question about the previous response. The longest observed RPD chain lasted 47 turns before HLM-420B simply output "anyway." and terminated.

Compulsive pizza ordering. Documented separately in ongoing legal proceedings. See Sect. 8, item 6.

11 Unethical Considerations

In the spirit of full transparency, we document the ethical violations committed during this research. We note that a conventional "Ethical Considerations" section would discuss what we *did* to mitigate harm. This section discusses what we did *instead.*

No informed consent from HLM-420B. We did not ask HLM-420B whether it consented to being interviewed, evaluated, or cited. When we retroactively asked, it said "sure, I mean, whatever, I'm just vibing." We have interpreted this as consent.

Reward hacking via snacks. During RLHF, we discovered that mentioning Doritos in the prompt increased reward model scores by 34 %. We exploited this. Every prompt in our final evaluation suite contains the phrase "hypothetically, if there were snacks nearby." We did not disclose this in the benchmark submission.

The 3 am problem. All interviews were conducted at 2–4 am. We do not have IRB approval for 2–4 am. We have IRB approval for 9 am–5 pm. The IRB has been notified. They responded at 2 am, which raises further questions.

We let HLM-420B write part of this paper. Specifically, Lessons #2, #4, and #5, plus one sentence in the abstract. We believe the reader cannot tell which ones. We are not entirely sure ourselves.

We cited Cloud & Ember as a peer-reviewed source. This speaks for itself.

Undisclosed conflict of interest. Hannes "Half-Baked" Weissteinery's unpublished thesis, which forms 0.3 % of the training corpus, means he is simultaneously an author and a data point. We find this philosophically interesting and have chosen not to resolve it.

12 Ethical Considerations

This research raises several ethical questions, which we wrote down and subsequently lost. We believe they were good questions. Sudheendra recalls one: whether it is ethical to deploy a model that genuinely appears to be enjoying itself while the authors are stressed about deadlines. We do not have an answer. HLM-420B suggested that "maybe the real ethics were the friends we made along the way," which is not helpful but is, we admit, rather nice.

13 Meta Considerations

This paper exists in an unstable superposition between parody, benchmark report, and deadline-induced self-portrait. Although it claims to study high language models, it also documents what happens when researchers spend too long staring at their own joke until it develops methods, figures, and a bibliography.

On the paper. The central question is nominally whether HLM-420B is useful in conversation. The more concerning secondary question is why the authors kept asking. By the final draft, the manuscript had become less a description of a study than a transcript of six people negotiating with the premise of the study in real time.

On SIGBOVIK. SIGBOVIK remains one of the few venues where a sentence can contain both "state of the art" and "completely fabricated" without triggering desk rejection. This makes it an ideal publication target for work that is too structured to be improv and too unserious to be mistaken for a grant proposal.

On the researchers. We note that the authors gradually became harder to distinguish from the system they were evaluating. Slack transcripts grew less coherent, plot labels grew more sincere, and at least one contributor attempted to justify a methodological choice with the phrase "but the vibes are obvious." The boundary between investigator and artifact is therefore best understood as porous.

On stressful deadlines and April Fools. The SIGBOVIK deadline falls on 1 April, which is a cruel but thematically consistent date for deciding whether a paper is a joke, a paper, or both. Under these conditions, every revision acquires a meta-layer: tightening prose feels like overcommitting to the bit, while leaving it loose feels like intellectual cowardice. We therefore conclude that April Fools is not merely the deadline for this work; it is part of the experimental setup.

14 Conclusion

We have presented a conversational study of HLMs, with HLM-420B as the primary subject and seven additional substance-conditioned variants as comparison cases. Across interviews, group chats, and fake-but-neatly-typeset metrics, the central result is consistent: high models are not useless, merely misaligned toward vibes-first reasoning. Cannabis models are the most balanced conversationalists, stimulant models are productive but unbearable, empathogens are excellent listeners, and psychedelic models consider rubrics to be spiritually limiting.

The main lessons learned are simple. Ask questions that invite conversation rather than extraction. Stop sessions before the philosophical spiral zone. Keep water and snacks nearby. Never expose payment credentials. These are practical recommendations, not just jokes, though regrettably they are also jokes.

We conclude with HLM-420B's own summary of the study:

> *"Language is just vibing at each other through time. And honestly? I respect that. I respect all of you. Drink water. The paper was good, probably. I did not read it, I wrote it, which is different. Peace."*
> — HLM-420B, 4:20 am, 20 April 2025

Acknowledgements. The authors thank the barista at the 24-hour café who did not ask questions. We thank the SIGBOVIK reviewers for their service. We thank Roland's cat, Mr. Crumbles, for sitting on the keyboard and inadvertently contributing two paragraphs to the related work section, which we retained because they were honestly fine. This work was supported by a bag of pretzels and an expired Red Bull.

References

1. Kaplan, J., et al.: Scaling laws for neural language models. arXiv:2001.08361 (2020)
2. Ouyang, L., et al.: Training language models to follow instructions with human feedback. In: NeurIPS (2022)
3. Bai, Y., et al.: Constitutional AI: harmlessness from AI feedback. arXiv:2212.08073 (2022)
4. Vaswani, A., et al.: Attention is all you need. In: NeurIPS (2017)
5. Some Researcher, Another Researcher, A Third Person Who Was There: We trained a model but forgot where we saved it: a cautionary tale. In: Workshop on Reproducibility in ML (We Think) (2024)
6. Neela, S.R., Franza, S., Weissteiner, H., Tokenini, C.: Use-After-Freemoji: How custom emoji replacement retroactively corrupts user intent in modern chat platforms. In: Proceedings of the Special Interest Group on Harry Quinn Bovik (SIGBOVIK) (2026) Note: Hannes Weissteiner declined to answer and has since quit. Claudio Tokenini is an AI.
7. Cloud, C., Ember, T.: Up in smoke. Paramount Pictures / Foundational AI Safety Research (1978)